

DATA-260 Homework 4, Poushali Purkayastha

Configuration

Value	Result
SID4	0571
PORT_BASE	$8000 + (571 \bmod 900) = 8571$
PREFIX	s0571
SEED	571
VERIFY_SEED	260571
DOMAIN_ID	$571 \bmod 8 = 3$, Grocery supply and recall notices
Hardware	Intel Core i5-1135G7 (4 cores, 8 threads), 15.8 GB RAM, Intel Iris Xe Graphics, Windows 11 Home, CPU only
Local model	qwen3:8b through Ollama (temperature 0, thinking off) for the RAG answers; sentence-transformers/all-MiniLM-L6-v2 (384-dimensional embeddings, CPU) for the chunk index
Database	MySQL 8.0.46 (Docker image mysql:8.0, container s0571-mysql on port 3307), database s0571_rel
Software	Python 3.12.10, fastapi 0.141.1, uvicorn 0.46.0, SQLAlchemy 2.0.52, PyMySQL 2.2.8, pydantic 2.13.3, Node 24.11.1, React 18, react-router-dom 6, Vite 5, axios 1, faiss-cpu 1.15.0, sentence-transformers 6.0.1, langchain-text-splitters 1.1.2
Repository	https://github.com/Poushali0202/data260-0571 (collaborators Sbnikitha and supriyaselvangesan)
Tagged commit	hw4 = 90386b46b4d22cb886f68805eb1ed9794bddf36d

The HW1 to HW3 code is extended in place. `app.py` keeps the notice API and the HW3 pages but the records now live in MySQL: `database.py` (engine and the session factory `db_session_basede26`), `models.py`, `schemas.py` and `routers/notices.py` are new, `routers/auth.py` was rewritten around a sessions table, and the React client lives in `frontend/`. Also new at the root: `seed_data.py`, `migrations/`, `measure_n_plus_one.py`, `rag.py`, `verify_hw04.py` (with the `verify-hw04` target in the Makefile), `.env.example`. Everything for this homework is in `reports/hw04/`: `RUN_LOG.txt`, `raw/`, `METRICS.md`, `AI_USE.md`, `README.md` (run instructions), `verification.json`, `report.pdf` and `screenshots/`.

Project folder structure (root, `frontend/` and `reports/hw04/`):

```
data260-0571/
├── app.py
├── database.py
├── models.py
├── schemas.py
├── seed_data.py
├── measure_n_plus_one.py
├── rag.py
├── verify_hw04.py
├── make_report_pdf.py
├── Makefile
├── requirements.txt
├── .env.example
├── README.md
├── routers/
│   ├── auth.py
│   ├── notices.py
│   ├── base.html
│   ├── index.html
│   ├── login.html
│   └── dashboard.html
├── templates/
├── migrations/
│   ├── hw04_schema.sql
│   └── hw04_add_index.sql
├── corpus/
│   └── 21 text documents (from HW3)
├── src/
│   └── model_client.py
├── frontend/
│   ├── index.html
│   ├── package.json
│   ├── vite.config.js
│   ├── src/
│   │   ├── main.jsx
│   │   ├── App.jsx
│   │   ├── api.js
│   │   └── styles.css
│   └── src/components/
│       └── Navbar.jsx
```

src/pages/
reports/hw04/
verification.json
raw/
rag_retrievals.txt
screenshots/

Login.jsx Home.jsx CreateRecord.jsx UpdateRecord.jsx DeleteRecord.jsx
RUN_LOG.txt METRICS.md AI_USE.md README.md report.md report.pdf
n_plus_one_requests.csv n_plus_one_summary.json rag_chunks.jsonl
rag_results.jsonl rag_comparison.md rag_k_sweep.md rag_evaluation.md

```
cap-project-structure
PS> Get-ChildItem (repository root) -Name
corpus corpus_warmup frontend
migrations reports routers
src templates .env
.env.example .gitignore 260_HW_1.pdf
AGENT.md agents_demo.py agent_graph.py
app.py chunking_compare.py common errors- docker-aws-wsl.pdf
database.py deploy_ecs.ps1 Docker and AWS ECS Deployment Guide.pdf
docker-compose.yml Dockerfile DOMAIN_SCHEMA.md
ecs-task-definition.json ecs-task-definition.rendered.json feedback.js
fetch_corpus.py hw1_client.py index.html
Makefile make_report_pdf.py measure_n_plus_one.py
models.py rag.py README.md
requirements.txt run_graph_experiment.py run_nondeterminism.py
schemas.py seed_data.py summarize_retrieval.py
verify_hw01.py verify_hw02.py verify_hw03.py
verify_hw04.py
PS> Get-ChildItem routers -Name
auth.py notices.py
PS> Get-ChildItem templates -Name
base.html dashboard.html index.html
login.html
PS> Get-ChildItem migrations -Name
hw04_add_index.sql hw04_schema.sql
PS> Get-ChildItem src -Name
model_client.py
PS> Get-ChildItem corpus -Name
ecfr_21_cfr_part_7_subpart_c_recalls.txt fda_101_product_recalls.txt fda_food_recalls_what_you_need_to_know.txt
fda_fsma_food_traceability_rule.txt fda_fsma_preventive_controls_human_food.txt fda_recalls_background_and_definitions.txt
openfda_food_enforcement_reports.txt usda_ers_food_service_market_segments.txt wikipedia_2006_spinach_e_coli_outbreak.txt
wikipedia_2008_canadian_listeriosis_outbreak.txt wikipedia_2011_germany_e_coli_outbreak.txt wikipedia_2022_infant_formula_shortage.txt
wikipedia_cold_chain.txt wikipedia_fda_food_safety_modernization_act.txt wikipedia_food_distribution.txt
wikipedia_food_safety_and_inspection_service.txt wikipedia_global_supply_chain_crisis.txt wikipedia_grocery_store.txt
wikipedia_peanut_corporation_of_america.txt wikipedia_product_recall.txt wikipedia_traceability.txt
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS-Docker_export>
```

```
cap-project-structure-2
PS> Get-ChildItem frontend -Name
src
package-lock.json
PS> Get-ChildItem frontend\src -Name
api.js
styles.css
PS> Get-ChildItem frontend\src\components -Name
Navbar.jsx
PS> Get-ChildItem frontend\src\pages -Name
CreateRecord.jsx
Login.jsx
PS> Get-ChildItem reports -Name
hw01
hw04
PS> Get-ChildItem reports\hw04 -Name
raw
METRICS.md
report.pdf
PS> Get-ChildItem reports\hw04\raw -Name
n_plus_one_requests.csv
rag_comparison.md
rag_results.jsonl
PS> (Get-ChildItem reports\hw04\screenshots).Count -> 43 screenshot files
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

Part 1: React client

The client is a Vite project in frontend/ (npm install, npm run dev, <http://localhost:5173>). It talks to the FastAPI backend on port 8571 through one axios instance with withCredentials: true, so the browser sends the HTTP-only session_id cookie with every request:

```
const api = axios.create({
  baseURL: "http://localhost:8571",
  withCredentials: true,
});

export async function login(email, password) {
  const res = await api.post("/auth/login", { email, password });
  return res.data;
}

export async function fetchNotices() {
  const res = await api.get("/api/notices");
  return res.data;
}
```

App.jsx owns the state (user, notices), checks the cookie session once on load with GET /auth/me, loads the list whenever a user is logged in, and passes the three handlers down as props. Every handler calls the backend, updates the state and redirects to the home page with useNavigate:

```
export default function App() {
  const navigate = useNavigate();
  const [user, setUser] = useState(null);
  const [notices, setNotices] = useState([]);
  useEffect(() => {
```

```

    me().then(setUser).catch(() => setUser(null));
  }, []);

  useEffect(() => {
    if (user) {
      fetchNotices().then(setNotices);
    } else {
      setNotices([]);
    }
  }, [user]);

  async function addNotice(data) {
    const created = await createNotice(data);
    setNotices([...notices, created]);
    navigate("/");
  }

  async function editNotice(id, data) {
    const updated = await updateNotice(id, data);
    setNotices(notices.map((n) => (n.id === id ? updated : n)));
    navigate("/");
  }

  async function removeNotice(id) {
    await deleteNotice(id);
    setNotices(notices.filter((n) => n.id !== id));
    navigate("/");
  }

  return (
    <div className="container">
      <Navbar user={user} onLogout={handleLogout} />
      <Routes>
        <Route path="/" element={<Home user={user} notices={notices} />} />
        <Route path="/login" element={<Login onLogin={setUser} />} />
        <Route path="/create" element={<CreateRecord user={user} onAdd={addNotice} />} />
        <Route path="/update/:id" element={<UpdateRecord user={user} onUpdate={editNotice} />} />
        <Route path="/delete/:id" element={<DeleteRecord user={user} onDelete={removeNotice} />} />
      </Routes>
    </div>
  );
}

```

Login page (Login.jsx)

The form posts email and password to POST /auth/login. On success the backend answers with the user and sets the cookie; the component stores the user through the onLogin prop and navigates to /. A wrong password shows the error text.

```

export default function Login({ onLogin }) {
  const navigate = useNavigate();
  const [email, setEmail] = useState("");
  const [password, setPassword] = useState("");
  const [error, setError] = useState("");

  async function handleSubmit(e) {
    e.preventDefault();
    try {
      const user = await login(email, password);
      onLogin(user);
      navigate("/");
    } catch {
      setError("Invalid email or password");
    }
  }

  ...
}

```

Grocery Notices

HomeLog in

Log in

Email

Password

Log in

Login page after a wrong password:

Grocery Notices

HomeLog in

Log in

Invalid email or password

Email

inspector@example.com

Password


.....

Log in

Only logged-in users see the list and the CRUD pages. `Navbar.jsx` hides the "Add notice" link while logged out, and `Home`, `CreateRecord`, `UpdateRecord` and `DeleteRecord` all return a "Login required" message when the `user` prop is null. The backend rejects the same routes with 401 anyway, so the guard is not only in the client.

```
{user ? (
  <>
    <NavLink to="/create">Add notice</NavLink>
    <span className="user">{user.name}</span>
    <button type="button" onClick={onLogout}>Log out</button>
  </>
) : (
  <NavLink to="/login">Log in</NavLink>
)}
```

Home page while logged out ("Login required", no Add button):



The cookie set by the login holds only the opaque token and is flagged `HttpOnly`: the browser context used for these screenshots reports `session_id` with `httpOnly=True` and `sameSite=Lax`, and the curl transcript in Part 2 shows the same `Set-Cookie` header next to the matching row of the sessions table.

1. Home page (`Home.jsx`, route `/`)

Renders every record in a table with an Update and a Delete link per row.

```
export default function Home({ user, notices }) {
  if (!user) {
```

```

return (
  <p className="notice">
    Login required. <Link to="/login">Log in</Link> to see the grocery notices.
  </p>
);
}

return (
  <div className="card">
    <div className="card-header">
      <h2>Grocery notices ({notices.length})</h2>
      <Link className="button" to="/create">Add notice</Link>
    </div>
    <table>
      <thead>
        <tr><th>ID</th><th>Product name</th><th>Source or manufacturer</th><th>Actions</th></tr>
      </thead>
      <tbody>
        {notices.map((notice) => (
          <tr key={notice.id}>
            <td>{notice.id}</td>
            <td>{notice.productName}</td>
            <td>{notice.noticeSource}</td>
            <td>
              <Link to={` /update/${notice.id}`}>Update</Link>
              <Link className="danger" to={` /delete/${notice.id}`}>Delete</Link>
            </td>
          </tr>
        ))}
      </tbody>
    </table>
  </div>
);
}

```

Grocery Notices			
		Home	Add notice Poushali Log out

Grocery notices (5000)			Add notice
ID	Product name	Source or manufacturer	Actions
1	Fresh chicken thighs 8 oz	Mission Organics	Update Delete
2	Store brand romaine lettuce 12 oz	USDA FSIS notice	Update Delete
3	Fresh tomato sauce 12 oz	Dairy Fresh Co-op	Update Delete
4	Bagged hummus 5 oz	Sunrise Bakery	Update Delete
5	Bagged tahini 8 oz	Natural Selection Foods	Update Delete
6	Organic eggs family pack	Sunrise Bakery	Update Delete
7	Bagged romaine lettuce 5 oz	Northern Meats	Update Delete
8	Store brand romaine lettuce 5 oz	Sunrise Bakery	Update Delete
9	Family size smoked salmon 1 lb	Evergreen Market	Update Delete
10	Fresh sliced mushrooms 12 oz	Summit Frozen Foods	Update Delete
11	Fresh whole milk family pack	Green Valley Farms	Update Delete
12	Canned whole milk family pack	Abbott Nutrition	Update Delete
13	Frozen deli ham family pack	Acme Foods	Update Delete
14	Frozen mixed berries 5 oz	Peanut Corporation of America	Update Delete

II. Add a new record (CreateRecord.jsx, route /create)

Two inputs (product name, the primary field; source or manufacturer, the secondary field) and an "Add notice" button. Submitting calls the onAdd prop; the backend inserts the row, MySQL assigns the auto-increment id, and App appends the returned record and redirects to /.

```
export default function CreateRecord({ user, onAdd }) {
  const [productName, setProductName] = useState("");
  const [noticeSource, setNoticeSource] = useState("");

  if (!user) {
    return <p className="notice">Login required</p>;
  }

  function handleSubmit(e) {
    e.preventDefault();
    onAdd({ productName, noticeSource });
  }

  return (
    <div className="card">
      <h2>Add a grocery notice</h2>
      <form onSubmit={handleSubmit}>
        <label>
          Product name
          <input value={productName} onChange={(e) => setProductName(e.target.value)} required />
        </label>
        <label>
          Source or manufacturer
          <input value={noticeSource} onChange={(e) => setNoticeSource(e.target.value)} required />
        </label>
        <button type="submit">Add notice</button>
      </form>
    </div>
  );
}
```

Add a grocery notice

Product name

Source or manufacturer

[Add notice](#)

Home page after the submit, with the new record and its id at the end of the list:

4983	Frozen cantaloupe 6 pack	Sunrise Bakery	Update	Delete
4984	Canned oat cereal 8 oz	Maple Leaf Foods	Update	Delete
4985	Store brand shredded coconut 8 oz	Northern Meats	Update	Delete
4986	Canned smoked salmon family pack	Peanut Corporation of America	Update	Delete
4987	Store brand almond flour 6 pack	Mission Organics	Update	Delete
4988	Organic peanut butter family pack	Golden State Growers	Update	Delete
4989	Organic eggs 5 oz	Northern Meats	Update	Delete
4990	Organic tomato sauce 5 oz	FDA recall bulletin	Update	Delete
4991	Organic oat cereal 8 oz	Abbott Nutrition	Update	Delete
4992	Family size tomato sauce 1 gallon	Acme Foods	Update	Delete
4993	Canned ice cream 1 lb	Riverbend Foods	Update	Delete
4994	Organic peanut butter 1 gallon	Riverbend Foods	Update	Delete
4995	Family size eggs 1 lb	Blue Ridge Creamery	Update	Delete
4996	Store brand eggs family pack	Coastal Distributors	Update	Delete
4997	Bagged cheddar cheese 1 lb	Evergreen Market	Update	Delete
4998	Frozen ground beef 5 oz	Abbott Nutrition	Update	Delete
4999	Fresh chicken thighs 12 oz	Harvest Fresh Produce	Update	Delete
5000	Dried tahini 1 gallon	Abbott Nutrition	Update	Delete
5003	Organic baby spinach 5 oz	FDA recall bulletin	Update	Delete

III. Update record (UpdateRecord.jsx, route /update/:id)

useParams gives the id, useEffect loads the current values from GET /api/notices/{id} into the two inputs, and the "Update notice" button calls the onUpdate prop, which sends PUT /api/notices/{id} and redirects home.

```
export default function UpdateRecord({ user, onUpdate }) {
  const { id } = useParams();
  const noticeId = Number(id);
  const [productName, setProductName] = useState("");
  const [noticeSource, setNoticeSource] = useState("");

  useEffect(() => {
    if (user) {
      fetchNotice(noticeId).then((notice) => {
        setProductName(notice.productName);
        setNoticeSource(notice.noticeSource);
      });
    }
  }, [user, noticeId]);

  if (!user) {
    return <p className="notice">Login required</p>;
  }

  function handleSubmit(e) {
    e.preventDefault();
    onUpdate(noticeId, { productName, noticeSource });
  }
  ...
}
```

Grocery Notices

HomeAdd noticePoushaliLog out

Update notice 5003

Product name

Organic baby spinach 5 oz bag

Source or manufacturer

FDA recall bulletin, updated

Update notice

Home page after the update:

4983	Frozen cantaloupe 6 pack	Sunrise Bakery	Update	Delete
4984	Canned oat cereal 8 oz	Maple Leaf Foods	Update	Delete
4985	Store brand shredded coconut 8 oz	Northern Meats	Update	Delete
4986	Canned smoked salmon family pack	Peanut Corporation of America	Update	Delete
4987	Store brand almond flour 6 pack	Mission Organics	Update	Delete
4988	Organic peanut butter family pack	Golden State Growers	Update	Delete
4989	Organic eggs 5 oz	Northern Meats	Update	Delete
4990	Organic tomato sauce 5 oz	FDA recall bulletin	Update	Delete
4991	Organic oat cereal 8 oz	Abbott Nutrition	Update	Delete
4992	Family size tomato sauce 1 gallon	Acme Foods	Update	Delete
4993	Canned ice cream 1 lb	Riverbend Foods	Update	Delete
4994	Organic peanut butter 1 gallon	Riverbend Foods	Update	Delete
4995	Family size eggs 1 lb	Blue Ridge Creamery	Update	Delete
4996	Store brand eggs family pack	Coastal Distributors	Update	Delete
4997	Bagged cheddar cheese 1 lb	Evergreen Market	Update	Delete
4998	Frozen ground beef 5 oz	Abbott Nutrition	Update	Delete
4999	Fresh chicken thighs 12 oz	Harvest Fresh Produce	Update	Delete
5000	Dried tahini 1 gallon	Abbott Nutrition	Update	Delete
5003	Organic baby spinach 5 oz bag	FDA recall bulletin, updated	Update	Delete

IV. Delete the record (DeleteRecord.jsx, route /delete/:id)

Loads the record to show what is about to be deleted; the "Delete notice" button calls the onDelete prop, which sends DELETE /api/notices/{id} and redirects home.


```
export default function DeleteRecord({ user, onDelete }) {
  const { id } = useParams();
  const noticeId = Number(id);
  const [notice, setNotice] = useState(null);

  useEffect(() => {
    if (user) {
      fetchNotice(noticeId).then(setNotice).catch(() => setNotice(null));
    }
  }, [user, noticeId]);

  if (!user) {
    return <p className="notice">Login required</p>;
  }

  return (
    <div className="card">
      <h2>Delete notice {noticeId}</h2>
      {notice ? (
        <>
          <p>Delete <strong>{notice.productName}</strong> from {notice.noticeSource}?</p>
          <button type="button" className="danger" onClick={() => onDelete(noticeId)}>
            Delete notice
          </button>
        </>
      ) : null}
    </div>
  );
}
```

```
    ) : (  
      <p className="notice">Notice not found.</p>  
    )}  
  </div>  
);  
}
```



Home page after the delete:

4982	Family size deli ham 1 lb	Acme Foods	Update	Delete
4983	Frozen cantaloupe 6 pack	Sunrise Bakery	Update	Delete
4984	Canned oat cereal 8 oz	Maple Leaf Foods	Update	Delete
4985	Store brand shredded coconut 8 oz	Northern Meats	Update	Delete
4986	Canned smoked salmon family pack	Peanut Corporation of America	Update	Delete
4987	Store brand almond flour 6 pack	Mission Organics	Update	Delete
4988	Organic peanut butter family pack	Golden State Growers	Update	Delete
4989	Organic eggs 5 oz	Northern Meats	Update	Delete
4990	Organic tomato sauce 5 oz	FDA recall bulletin	Update	Delete
4991	Organic oat cereal 8 oz	Abbott Nutrition	Update	Delete
4992	Family size tomato sauce 1 gallon	Acme Foods	Update	Delete
4993	Canned ice cream 1 lb	Riverbend Foods	Update	Delete
4994	Organic peanut butter 1 gallon	Riverbend Foods	Update	Delete
4995	Family size eggs 1 lb	Blue Ridge Creamery	Update	Delete
4996	Store brand eggs family pack	Coastal Distributors	Update	Delete
4997	Bagged cheddar cheese 1 lb	Evergreen Market	Update	Delete
4998	Frozen ground beef 5 oz	Abbott Nutrition	Update	Delete
4999	Fresh chicken thighs 12 oz	Harvest Fresh Produce	Update	Delete
5000	Dried tahini 1 gallon	Abbott Nutrition	Update	Delete

V. Props, hooks and routing

Create, Update and Delete receive their handler (onAdd, onUpdate, onDelete) and the user as props from App; Home receives user and notices; Login receives onLogin. useState holds the form values and the loaded record, useEffect runs the session check, the list load and the record load for the update and delete pages. Routing is react-router-dom 6: BrowserRouter in main.jsx, Routes/Route in App.jsx, Link/NavLink in the pages, useParams for the id and useNavigate for the redirects.

Part 2: MySQL persistence and server-side sessions

Database

Database s0571_rel (PREFIX s0571 plus _rel) with four tables. seed_data.py creates the database if needed and the tables from the SQLAlchemy models; the resulting DDL is committed as migrations/hw04_schema.sql:

```
CREATE DATABASE IF NOT EXISTS s0571_rel;
USE s0571_rel;

CREATE TABLE `grocery_notices` (
  `id` int NOT NULL AUTO_INCREMENT,
  `product_name` varchar(200) NOT NULL,
  `notice_source` varchar(200) NOT NULL,
  PRIMARY KEY (`id`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci;
```

```

CREATE TABLE `notice_lots` (
  `id` int NOT NULL AUTO_INCREMENT,
  `notice_id` int NOT NULL,
  `lot_code` varchar(50) NOT NULL,
  `units_affected` int NOT NULL,
  PRIMARY KEY (`id`),
  KEY `notice_id` (`notice_id`),
  CONSTRAINT `notice_lots_ibfk_1` FOREIGN KEY (`notice_id`) REFERENCES `grocery_notices` (`id`) ON
DELETE CASCADE
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci;
CREATE TABLE `sessions` (
  `id` varchar(64) NOT NULL,
  `user_id` int NOT NULL,
  `created_at` datetime NOT NULL,
  `expires_at` datetime NOT NULL,
  PRIMARY KEY (`id`),
  KEY `user_id` (`user_id`),
  CONSTRAINT `sessions_ibfk_1` FOREIGN KEY (`user_id`) REFERENCES `users` (`id`) ON DELETE CASCADE
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci;
CREATE TABLE `users` (
  `id` int NOT NULL AUTO_INCREMENT,
  `name` varchar(100) NOT NULL,
  `email` varchar(255) NOT NULL,
  `password_hash` varchar(255) NOT NULL,
  PRIMARY KEY (`id`),
  UNIQUE KEY `email` (`email`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci;

```

grocery_notices is the primary domain entity (auto-increment id, primary field product_name, secondary field notice_source); notice_lots is the related test data for Part 3; users and sessions are the authentication tables asked for.

database.py reads DATABASE_URL from .env (.env.example shows the format) and defines the connection variable db_session_based26:

```

load_dotenv()

DATABASE_URL = os.getenv("DATABASE_URL")
if not DATABASE_URL:
    raise RuntimeError("DATABASE_URL is missing. Copy .env.example to .env and fill in the MySQL
password.")

engine = create_engine(DATABASE_URL, pool_pre_ping=True)
db_session_based26 = sessionmaker(bind=engine, autoflush=False, expire_on_commit=False)
Base = declarative_base()

def get_db():
    db = db_session_based26()
    try:
        yield db
    finally:
        db.close()

```

models.py:

```

class GroceryNotice(Base):
    __tablename__ = "grocery_notices"

    id = Column(Integer, primary_key=True, autoincrement=True)
    productName = Column("product_name", String(200), nullable=False)
    noticeSource = Column("notice_source", String(200), nullable=False)
    lots = relationship("NoticeLot", cascade="all, delete-orphan")

class NoticeLot(Base):
    __tablename__ = "notice_lots"
    id = Column(Integer, primary_key=True, autoincrement=True)

```

```

        notice_id = Column(Integer, ForeignKey("grocery_notices.id", ondelete="CASCADE"),
nullable=False)
        lotCode = Column("lot_code", String(50), nullable=False)
        unitsAffected = Column("units_affected", Integer, nullable=False)

class User(Base):
    __tablename__ = "users"

    id = Column(Integer, primary_key=True, autoincrement=True)
    name = Column(String(100), nullable=False)
    email = Column(String(255), nullable=False, unique=True)
    password_hash = Column(String(255), nullable=False)

class UserSession(Base):
    __tablename__ = "sessions"

    id = Column(String(64), primary_key=True)
    user_id = Column(Integer, ForeignKey("users.id", ondelete="CASCADE"), nullable=False)
    created_at = Column(DateTime, nullable=False, default=datetime.now)
    expires_at = Column(DateTime, nullable=False)
    user = relationship("User", lazy="joined")

```

Database s0571_re1 seen through the mysql client: the four tables with their columns, then the row counts, sample rows, the two users and the sessions table. Postman and Workbench are not installed on this laptop, so the client screenshots in this report come from the mysql command line and from the API client built into FastAPI (Swagger UI at /docs).

```
cap-db-schema
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "SHOW TABLES; DESCRIBE grocery_notices; DESCRIBE notice
_lots; DESCRIBE users"
+-----+
| Tables_in_s0571_rel |
+-----+
| grocery_notices
| notice_lots
| sessions
| users
+-----+

+-----+-----+-----+-----+-----+-----+
| Field      | Type      | Null | Key | Default | Extra      |
+-----+-----+-----+-----+-----+-----+
| id          | int       | NO   | PRI | NULL    | auto_increment
| product_name | varchar(200) | NO   |     | NULL    |
| notice_source | varchar(200) | NO   | MUL | NULL    |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+-----+
| Field      | Type      | Null | Key | Default | Extra      |
+-----+-----+-----+-----+-----+-----+
| id          | int       | NO   | PRI | NULL    | auto_increment
| notice_id   | int       | NO   | MUL | NULL    |
| lot_code    | varchar(50) | NO   |     | NULL    |
| units_affected | int       | NO   |     | NULL    |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+-----+
| Field      | Type      | Null | Key | Default | Extra      |
+-----+-----+-----+-----+-----+-----+
| id          | int       | NO   | PRI | NULL    | auto_increment
| name        | varchar(100) | NO   |     | NULL    |
| email       | varchar(255) | NO   | UNI | NULL    |
| password_hash | varchar(255) | NO   |     | NULL    |
+-----+-----+-----+-----+-----+-----+
PS C:\Users\Boushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

```
cap-db-rows
```

Field	Type	Null	Key	Default	Extra
id	varchar(64)	NO	PRI	NULL	
user_id	int	NO	MUL	NULL	
created_at	datetime	NO		NULL	
expires_at	datetime	NO		NULL	

```
PS> mysql ... s0571_rel -e "SELECT (SELECT COUNT(*) FROM grocery_notices) AS notices, (SELECT COUNT(*) FROM notice_lots) AS lots, (SELECT COUNT(*) FROM users) AS users, (SELECT COUNT(*) FROM sessions) AS sessions"
```

notices	lots	users	sessions
5000	200	2	1

```
PS> mysql ... s0571_rel -e "SELECT * FROM grocery_notices ORDER BY id DESC LIMIT 3; SELECT * FROM notice_lots LIMIT 3"
```

id	product_name	notice_source
5000	Dried tahini 1 gallon	Abbott Nutrition
4999	Fresh chicken thighs 12 oz	Harvest Fresh Produce
4998	Frozen ground beef 5 oz	Abbott Nutrition

id	notice_id	lot_code	units_affected
1	4896	LOT-977337	3043
2	4889	LOT-413239	221
3	4991	LOT-926251	4635

```
PS> mysql ... s0571_rel -e "SELECT id, name, email, LEFT(password_hash, 24) AS password_hash FROM users; SELECT LEFT(id, 20) AS session_token, user_id, created_at, expires_at FROM sessions"
```

id	name	email	password_hash
1	Poushali	inspector@example.com	e8dd057e305674d75af3aa41
2	Store Manager	manager@example.com	4f71040a05657006ec7191c6

session_token	user_id	created_at	expires_at
260bc2364ab2a99e771d	1	2026-09-27 23:08:24	2026-09-27 23:38:24

```
$ mysql ... -e "SHOW TABLES; SELECT COUNT(*) AS notices FROM grocery_notices; SELECT COUNT(*) AS lots FROM notice_lots; SELECT COUNT(*) AS users FROM users; ..."
```

```
Tables_in_s0571_rel
```

```
grocery_notices
```

```
notice_lots
```

```
sessions
```

```
users
```

```
notices
```

```
5000
```

```
lots
```

```
200
```

```
users
```

```
2
```

```
id      product_name      notice_source
5000    Dried tahini 1 gallon  Abbott Nutrition
4999    Fresh chicken thighs 12 oz  Harvest Fresh Produce
4998    Frozen ground beef 5 oz  Abbott Nutrition
```

```
id      notice_id  lot_code  units_affected
1        4896      LOT-977337  3043
2        4889      LOT-413239  221
3        4991      LOT-926251  4635
```

```
id      name      email
1        Poushali  inspector@example.com
2        Store Manager  manager@example.com
```

Server-side sessions

Passwords are stored as PBKDF2-SHA256 hashes with a random salt. A successful login inserts a row into sessions whose id is a random 64-character token and sets that token, and nothing else, in an HTTP-only cookie. Every protected route reads the cookie, looks the token up in the table, checks expires_at, and gets the user from the row; logout deletes the row and clears the cookie, so a copied cookie is useless afterwards (routers/auth.py):

```
def open_session(db: Session, user: User, response: Response):
    token = secrets.token_hex(32)
    now = datetime.now()
    expires = now + timedelta(minutes=SESSION_MINUTES)
    db.add(UserSession(id=token, user_id=user.id, created_at=now, expires_at=expires))
    db.commit()
    response.set_cookie(COOKIE, token, httponly=True, samesite="lax", max_age=SESSION_MINUTES * 60)

def close_session(db: Session, request: Request, response: Response):
    token = request.cookies.get(COOKIE)
    if token:
        db.query(UserSession).filter(UserSession.id == token).delete()
        db.commit()
    response.delete_cookie(COOKIE)

def current_user(request: Request, db: Session = Depends(get_db)):
    token = request.cookies.get(COOKIE)
    if not token:
        return None
    session = db.get(UserSession, token)
    if session is None:
        return None
    if session.expires_at < datetime.now():
        db.delete(session)
        db.commit()
        return None
    return session.user

def require_login(user=Depends(current_user)):
    if user is None:
        raise HTTPException(status_code=401, detail="Login required")
    return user

@router.post("/auth/login", response_model=UserOut)
def login(data: LoginIn, response: Response, db: Session = Depends(get_db)):
    user = find_user(db, data.email, data.password)
    if user is None:
        raise HTTPException(status_code=401, detail="Invalid email or password")
    open_session(db, user, response)
    return user
```

The curl transcript from RUN_LOG.txt shows the cookie and the matching row in the table:

```
$ curl -i -s -c cookies.txt -X POST -H "Content-Type: application/json" -d '{"email": "\
"inspector@example.com", "password": "\recall2026"}' http://127.0.0.1:8571/auth/login
HTTP/1.1 200 OK
set-cookie: session_id=38d14aac31fc093848d08f915610bc07a5ad1a11019e76868e5db6765fde92e8; HttpOnly;
Max-Age=1800; Path=/; SameSite=lax

{"id":1,"name":"Poushali","email":"inspector@example.com"}

$ docker exec s0571-mysql mysql -uroot -p... s0571_rel -e "SELECT id, user_id, created_at,
expires_at FROM sessions"
id                                     user_id  created_at
expires_at
38d14aac31fc093848d08f915610bc07a5ad1a11019e76868e5db6765fde92e8      1          2026-09-23 18:21:08
2026-09-23 18:51:08
$ curl -i -s -b cookies.txt -X POST http://127.0.0.1:8571/auth/logout
```

```

HTTP/1.1 200 OK
set-cookie: session_id=""; expires=Thu, 24 Sep 2026 01:21:09 GMT; Max-Age=0; Path=/; SameSite=lax

$ curl -i -s -b cookies.txt http://127.0.0.1:8571/auth/me    (same cookie after logout)
HTTP/1.1 401 Unauthorized
{"detail":"Login required"}

```

CRUD endpoints

All five operations are in `routers/notices.py`; the router-level dependency makes every one of them require the cookie session.

```

router = APIRouter(prefix="/api/notices", dependencies=[Depends(require_login)])

def find_notice(db: Session, notice_id: int) -> GroceryNotice:
    notice = db.get(GroceryNotice, notice_id)
    if notice is None:
        raise HTTPException(status_code=404, detail="Notice not found")
    return notice

@router.get("", response_model=list[NoticeOut])
def list_notices(q: str = "", db: Session = Depends(get_db)):
    query = db.query(GroceryNotice)
    if q.strip():
        term = q.strip()
        query = query.filter(GroceryNotice.productName.contains(term) | GroceryNotice.noticeSource.contains(term))
    return query.order_by(GroceryNotice.id).all()

@router.get("/{notice_id}", response_model=NoticeOut)
def get_notice(notice_id: int, db: Session = Depends(get_db)):
    return find_notice(db, notice_id)

@router.post("", response_model=NoticeOut, status_code=201)
def create_notice(data: NoticeIn, db: Session = Depends(get_db)):
    notice = GroceryNotice(**data.model_dump())
    db.add(notice)
    db.commit()
    db.refresh(notice)
    return notice

@router.put("/{notice_id}", response_model=NoticeOut)
def update_notice(notice_id: int, data: NoticeIn, db: Session = Depends(get_db)):
    notice = find_notice(db, notice_id)
    notice.productName = data.productName
    notice.noticeSource = data.noticeSource
    db.commit()
    db.refresh(notice)
    return notice

@router.delete("/{notice_id}", status_code=204)
def delete_notice(notice_id: int, db: Session = Depends(get_db)):
    db.delete(find_notice(db, notice_id))
    db.commit()

```

The requests below were sent from the API client built into FastAPI (Swagger UI at `http://localhost:8571/docs`), which keeps the session cookie between calls the way Postman does. Each screenshot shows the request, its curl equivalent, the status code and the response body.

Operation	Method and path	Screenshot
Without the cookie	GET /api/notices	api-unauthorized.png (401)
Log in (sets the cookie)	POST /auth/login	api-login.png
Add a new record	POST /api/notices	api-create.png (201)
View a record by id	GET /api/notices/{id}	api-get-by-id.png
Update record details	PUT /api/notices/{id}	api-update.png
Delete a record	DELETE /api/notices/{id}	api-delete.png (204)
View all records	GET /api/notices	api-list.png
Log out	POST /auth/logout	api-logout.png

GET /api/notices List Notices

Parameters
Cancel

Name	Description
q	q

Execute
Clear

Responses

Curl

```
curl -X 'GET' \
  'http://localhost:8571/api/notices' \
  -H 'accept: */*'

```

Request URL

```
http://localhost:8571/api/notices

```

Server response

Code	Details
401 <i>Undocumented</i>	Error: Unauthorized Response body <pre>{ "detail": "Login required" }</pre> Download Response headers <pre>content-length: 27 content-type: application/json date: Mon, 28 Sep 2026 05:56:50 GMT server: uvicorn </pre>

Responses

Code	Description	Links
200	Successful Response Media type application/json Controls Accept header. Example Value Schema <pre>[{ "noticeName": "string" }]</pre>	No links

POST /auth/login Login

Parameters

Cancel

Reset

No parameters

Request body required

application/json

Edit Value | Schema

```
{
  "email": "inspector@example.com",
  "password": "recall2026"
}
```

Execute

Clear

Responses

Curl

```
curl -X 'POST' \
'http://localhost:8571/auth/login' \
-H 'accept: */*' \
-H 'Content-Type: application/json' \
-d '{
  "email": "inspector@example.com",
  "password": "recall2026"
}'
```

Request URL

```
http://localhost:8571/auth/login
```

Server response

Code	Details
200	Response body <pre>{ "id": 1, "name": "Poushali", "email": "inspector@example.com" }</pre> Download

POST

/api/notices

Create Notice

^

Parameters

Cancel

Reset

No parameters

Request body

required

application/json

Edit Value

Schema

```
{
  "productName": "Frozen mixed berries 12 oz",
  "noticeSource": "Acme Foods"
}
```

Execute

Clear

Responses

Curl

```
curl -X 'POST' \
  'http://localhost:8571/api/notices' \
  -H 'accept: */*' \
  -H 'Content-Type: application/json' \
  -d '{
    "productName": "Frozen mixed berries 12 oz",
    "noticeSource": "Acme Foods"
  }'
```

Request URL

http://localhost:8571/api/notices

Server response

Code	Details
201	Response body <pre>{ "productName": "Frozen mixed berries 12 oz", "noticeSource": "Acme Foods", "id": 5007 }</pre> Download

GET

/api/notices/{notice_id}

Get Notice

⌵

Parameters

Cancel

Name	Description
notice_id * required	
integer	
(path)	

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
'http://localhost:8571/api/notices/5007' \
-H 'accept: */*'

```

Request URL

```
http://localhost:8571/api/notices/5007

```

Server response

Code	Details
200	Response body <pre>{ "productName": "Frozen mixed berries 12 oz", "noticeSource": "Acme Foods", "id": 5007 } </pre> Download Response headers <pre>content-length: 82 content-type: application/json date: Mon, 28 Sep 2026 05:57:00 GMT server: unicorn </pre>

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

Controls Accept header.

Example Value | Schema

```
{
  "productName": "string",
  "noticeSource": "string",

```

PUT

/api/notices/{notice_id}

Update Notice

^

Parameters

Cancel

Reset

Name	Description
notice_id * required	
integer	5007
(path)	

Request body required

application/json

Edit Value | Schema

```
{
  "productName": "Frozen mixed berries 12 oz bag",
  "noticeSource": "Acme Foods, updated"
}
```

Execute

Clear

Responses

Curl

```
curl -X 'PUT' \
'http://localhost:8571/api/notices/5007' \
-H 'accept: */*' \
-H 'Content-Type: application/json' \
-d '{
  "productName": "Frozen mixed berries 12 oz bag",
  "noticeSource": "Acme Foods, updated"
}'
```

Request URL

```
http://localhost:8571/api/notices/5007
```

Server response

Code	Details
200	Response body

DELETE

/api/notices/{notice_id}

Delete Notice

^

Parameters

Cancel

Name	Description
notice_id * required	
integer	5007
(path)	

Execute

Clear

Responses

Curl

```
curl -X 'DELETE' \
'http://localhost:8571/api/notices/5007' \
-H 'accept: */*'
```

Request URL

```
http://localhost:8571/api/notices/5007
```

Server response

Code	Details
204	Response headers access-control-allow-credentials: truecontent-type: application/jsondate: Mon, 28 Sep 2026 05:57:06 GMTserver: uvicorn

Responses

Code	Description	Links
204	Successful Response	No links
422	Validation Error	No links

Media type

application/json

Example Value | Schema

```
{  "detail": [    {      "loc": [        "string",        0      ],      "msg": "string",      "type": "string",      "input": "string",    }  ]}
```

GET

/api/notices

List Notices

^

Parameters

Cancel

Name	Description
q	
string	q
(query)	

ExecuteClear

Responses

Curl

curl -X 'GET' \ 'http://localhost:8571/api/notices' \ -H 'accept: */*' 

Request URL

http://localhost:8571/api/notices

Server response

Code	Details
200	Response body <pre>[{ "productName": "Fresh chicken thighs 8 oz", "noticeSource": "Mission Organics", "id": 1 }, { "productName": "Store brand romaine lettuce 12 oz", "noticeSource": "USDA FSIS notice", "id": 2 }, { "productName": "Fresh tomato sauce 12 oz", "noticeSource": "Dairy Fresh Co-op", "id": 3 }, { "productName": "Bagged hummus 5 oz", "noticeSource": "Sunrise Bakery", "id": 4 }, { "productName": "Bagged tahini 8 oz", "noticeSource": "Natural Selection Foods", "id": 5 }, { "productName": "Organic eggs family pack", </pre> Download Response headers content-length: 454383 content-type: application/json

POST

/auth/logout Logout

Parameters

No parameters

Execute

Clear

Responses

Curl

```
curl -X 'POST' \
'http://localhost:8571/auth/logout' \
-H 'accept: */*' \
-d ''
```

Request URL

http://localhost:8571/auth/logout

Server response

Code	Details
200	Response body <pre>{ "message": "logged out" }</pre> Response headers <pre>access-control-allow-credentials: true content-length: 24 content-type: application/json date: Mon, 28 Sep 2026 05:58:14 GMT server: unicorn</pre>

Responses

Code	Description	Links
200	Successful Response Media type application/json Controls Accept header. Example Value Schema <pre>"string"</pre>	No links

GET

/auth/me Me

The same round trip recorded with curl (RUN_LOG.txt):

```
$ curl -i -s -b cookies.txt -X POST -H "Content-Type: application/json" -d '{"productName":"Organic baby spinach 5 oz","noticeSource":"FDA recall bulletin"}' http://127.0.0.1:8571/api/notices
HTTP/1.1 201 Created
{"productName":"Organic baby spinach 5 oz","noticeSource":"FDA recall bulletin","id":5001}

$ curl -s -b cookies.txt http://127.0.0.1:8571/api/notices/5001
{"productName":"Organic baby spinach 5 oz","noticeSource":"FDA recall bulletin","id":5001}

$ curl -s -b cookies.txt -X PUT -H "Content-Type: application/json" -d '{"productName":"Organic baby spinach 5 oz bag","noticeSource":"FDA recall bulletin, updated"}' http://127.0.0.1:8571/api/notices/5001
{"productName":"Organic baby spinach 5 oz bag","noticeSource":"FDA recall bulletin, updated","id":5001}

$ curl -i -s -b cookies.txt -X DELETE http://127.0.0.1:8571/api/notices/5001
```

HTTP/1.1 204 No Content

\$ curl -i -s -b cookies.txt http://127.0.0.1:8571/api/notices/5001

HTTP/1.1 404 Not Found

{"detail":"Notice not found"}

```
cap-curl-checks
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part 2: API checks with curl)
=== Part 2: API checks with curl (server started with python app.py) ===
$ date
2026-09-23T18:21:07:00
$ curl -i -s http://127.0.0.1:8571/api/notices (no cookie)
HTTP/1.1 401 Unauthorized
date: Thu, 24 Sep 2026 01:21:07 GMT
server: uvicorn
content-length: 27
content-type: application/json

{"detail":"Login required"}

$ curl -i -s -X POST -H "Content-Type: application/json" -d '{"email":"inspector@example.com","password":"wrong"}' http://127.0.0.1:8571/auth/login
HTTP/1.1 401 Unauthorized
date: Thu, 24 Sep 2026 01:21:07 GMT
server: uvicorn
content-length: 38
content-type: application/json

{"detail":"Invalid email or password"}

$ curl -i -s -c cookies.txt -X POST -H "Content-Type: application/json" -d '{"email":"inspector@example.com","password":"recall2026"}' http://127.0.0.1:8571/auth/login
HTTP/1.1 200 OK
date: Thu, 24 Sep 2026 01:21:07 GMT
server: uvicorn
content-length: 58
content-type: application/json
set-cookie: session_id=38d14aac31fc093848d08f915610bc07a5ad1a11019e76868e5db6765fde92e8; HttpOnly; Max-Age=1800; Path=/; SameSite=lax

{"id":1,"name":"Poushali","email":"inspector@example.com"}

$ curl -s -b cookies.txt http://127.0.0.1:8571/auth/me
{"id":1,"name":"Poushali","email":"inspector@example.com"}
```

```
cap-curl-checks-2
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part 2: API checks with curl)
$ docker exec s0571-mysql mysql -uroot -p... s0571_rel -e "SELECT id, user_id, created_at, expires_at FROM sessions"
id      user_id created_at      expires_at
38d14aac31fc093848d08f915610bc07a5ad1a11019e76868e5db6765fde92e8      1      2026-09-23 18:21:08      2026-09-23 18:51:08

$ curl -i -s -b cookies.txt -X POST -H "Content-Type: application/json" -d '{"productName":"Organic baby spinach 5 oz","noticeSource":"FDA recall bulletin"}' http://127.0.0.1:8571/api/notices
HTTP/1.1 201 Created
date: Thu, 24 Sep 2026 01:21:07 GMT
server: uvicorn
content-length: 90
content-type: application/json

{"productName":"Organic baby spinach 5 oz","noticeSource":"FDA recall bulletin","id":5001}

$ curl -s -b cookies.txt http://127.0.0.1:8571/api/notices/5001
{"productName":"Organic baby spinach 5 oz","noticeSource":"FDA recall bulletin","id":5001}

$ curl -s -b cookies.txt -X PUT -H "Content-Type: application/json" -d '{"productName":"Organic baby spinach 5 oz bag","noticeSource":"FDA recall bulletin, updated"}' http://127.0.0.1:8571/api/notices/5001
{"productName":"Organic baby spinach 5 oz bag","noticeSource":"FDA recall bulletin, updated","id":5001}

$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices?q=spinach" | python -c "import json,sys; rows=json.load(sys.stdin); print(len(rows), \"rows\", last:\\\", rows[-1])"
221 rows, last: {'productName': 'Organic baby spinach 5 oz bag', 'noticeSource': 'FDA recall bulletin, updated', 'id': 5001}

$ curl -i -s -b cookies.txt -X DELETE http://127.0.0.1:8571/api/notices/5001
HTTP/1.1 204 No Content
date: Thu, 24 Sep 2026 01:21:08 GMT
server: uvicorn
content-type: application/json

$ curl -i -s -b cookies.txt http://127.0.0.1:8571/api/notices/5001
HTTP/1.1 404 Not Found
date: Thu, 24 Sep 2026 01:21:08 GMT
server: uvicorn
content-length: 29
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

```
cap-curl-checks-3
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part 2: API checks with curl)
content-type: application/json

{"detail":"Notice not found"}

$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices/naive?page_size=10" | python -c "import json,sys; b=json.load(sys.stdin); print({k: v for k, v in b.items() if k != \"notices\"}, \"first notice:\\\", b[\"notices\"][0])"
{'version': 'naive', 'page_size': 10, 'sql_queries': 11} first notice: {'productName': 'Dried tahini 1 gallon', 'noticeSource': 'Abbott Nutrition', 'id': 5000, 'lots': [{'id': 183, 'lotCode': 'LOT-553350', 'unitsAffected': 2298}]}

$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices/fixe?page_size=10" | python -c "..."
{'version': 'fixed', 'page_size': 10, 'sql_queries': 2} first notice: {'productName': 'Dried tahini 1 gallon', 'noticeSource': 'Abbott Nutrition', 'id': 5000, 'lots': [{'id': 183, 'lotCode': 'LOT-553350', 'unitsAffected': 2298}]}

$ curl -i -s -b cookies.txt -X POST http://127.0.0.1:8571/auth/logout
HTTP/1.1 200 OK
date: Thu, 24 Sep 2026 01:21:08 GMT
server: uvicorn
content-length: 24
content-type: application/json
set-cookie: session_id=""; expires=Thu, 24 Sep 2026 01:21:09 GMT; Max-Age=0; Path=/; SameSite=lax

{"message":"logged out"}

$ curl -i -s -b cookies.txt http://127.0.0.1:8571/auth/me (same cookie after logout)
HTTP/1.1 401 Unauthorized
date: Thu, 24 Sep 2026 01:21:08 GMT
server: uvicorn
content-length: 27
content-type: application/json

{"detail":"Login required"}

$ date
```

Part 3: N+1 measurement and query tuning

1 and 2. Seed data and the related table

seed_data.py seeds random with SEED 571, drops and recreates the tables, inserts the two accounts, 5,000 notices built from word lists (form, product, size, source) and 200 lots. Each lot belongs to one of the 200 newest notices, so the pages returned by the list endpoints (newest first) carry related data: a page of 10 has 9 lots, a page of 50 has 56, a page of 200 has all 200.

```
SEED = 571
NOTICE_COUNT = 5000
LOT_COUNT = 200

def main():
    random.seed(SEED)
    create_database()
    Base.metadata.drop_all(bind=engine)
    Base.metadata.create_all(bind=engine)

    db = db_session_based26()
    db.add_all([User(name=name, email=email, password_hash=hash_password(password)) for name, email,
password in USERS])
    notices = []
    for _ in range(NOTICE_COUNT):
        product = f"{random.choice(FORMS)} {random.choice(PRODUCTS)} {random.choice(SIZES)}"
        notices.append(GroceryNotice(productName=product, noticeSource=random.choice(SOURCES)))
    db.add_all(notices)
    db.commit()

    lots = []
    for _ in range(LOT_COUNT):
        lots.append(NoticeLot(
            notice_id=random.randint(NOTICE_COUNT - 199, NOTICE_COUNT),
            lotCode=f"LOT-{random.randint(100000, 999999)}",
            unitsAffected=random.randint(10, 5000),
        ))
    db.add_all(lots)
    db.commit()

$ date
2026-09-23T18:18:37-07:00
$ python seed_data.py
seed 571: 2 users, 5000 notices, 200 lots
$ date
2026-09-23T18:18:57-07:00
```

```
cap-seed x + v
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (sec
tion: Part 2 and 3: seeding)
=== Part 2 and 3: seeding (python seed_data.py) ===
$ date
2026-09-23T18:18:37-07:00
$ python seed_data.py
seed 571: 2 users, 5000 notices, 200 lots
$ date
2026-09-23T18:18:57-07:00
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

3 and 5. The naive endpoint and the fixed endpoint

Both endpoints return the same page (page_size newest notices, each with its lots) and the number of SQL statements the endpoint body issued. The count comes from a before_cursor_execute listener on the engine that increments a per-request counter kept in a ContextVar (database.py):

```
query_counter = ContextVar("query_counter", default=None)

@event.listens_for(engine, "before_cursor_execute")
def count_query(conn, cursor, statement, parameters, context, executemany):
    counter = query_counter.get()
    if counter is not None:
        counter[0] += 1

def start_counting():
    query_counter.set([0])

def queries_so_far():
    return query_counter.get()[0]
```

The naive version loads the notices and then lets NoticeWithLots.model_validate read notice.lots for each one; the relationship is lazy, so every access is one more SELECT ... FROM notice_lots WHERE notice_id = %s. The fixed version adds selectinload, which loads all lots of the page in one WHERE notice_id IN (...) statement:

```
@router.get("/naive", response_model=NoticePage)
def list_naive(page_size: int = 10, db: Session = Depends(get_db)):
    start_counting()
    notices = db.query(GroceryNotice).order_by(GroceryNotice.id.desc()).limit(page_size).all()
    items = [NoticeWithLots.model_validate(notice) for notice in notices]
    return NoticePage(version="naive", page_size=page_size, sql_queries=queries_so_far(),
notices=items)
@router.get("/fixed", response_model=NoticePage)
```

```

def list_fixed(page_size: int = 10, db: Session = Depends(get_db)):
    start_counting()
    notices = (
        db.query(GroceryNotice)
        .options(selectinload(GroceryNotice.lots))
        .order_by(GroceryNotice.id.desc())
        .limit(page_size)
        .all()
    )
    items = [NoticeWithLots.model_validate(notice) for notice in notices]
    return NoticePage(version="fixed", page_size=page_size, sql_queries=queries_so_far(),
    notices=items)

```

The MySQL general log confirms the counts for one request of each version at page size 10 (RUN_LOG.txt): the naive request produced one SELECT ... grocery_notices and ten SELECT ... notice_lots ... WHERE notice_id = %, the fixed request one notices query and one notice_lots ... IN (...) query. Every protected route also runs one SELECT sessions for the login check before the endpoint body; that statement is outside the reported count.

```

$ mysql ... -e "SELECT LEFT(argument, 22) AS statement_start, COUNT(*) AS statements FROM mysql.
general_log WHERE ... GROUP BY 1"
statement_start      statements
SELECT users.id AS use  1
SELECT sessions.id AS  2
SELECT grocery_notices 2
SELECT notice_lots.id  10
SELECT notice_lots.not  1

```

```

cap-general-log
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part 3 check:
statement counts confirmed with the MySQL general log ==)
=== Part 3 check: statement counts confirmed with the MySQL general log ===
$ date
2026-09-23T18:24:10-07:00
$ mysql ... -e "SET GLOBAL log_output = 'TABLE'; SET GLOBAL general_log = 'ON';"
$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices/naive?page_size=10" > /dev/null
$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices/fixed?page_size=10" > /dev/null
$ mysql ... -e "SET GLOBAL general_log = 'OFF';"
$ mysql ... -e "SELECT event_time, LEFT(argument, 95) AS statement FROM mysql.general_log WHERE command_type = 'Query' AND argument LIKE 'SE
LECT %' AND argument NOT LIKE '%general_log%' ORDER BY event_time"
event_time      statement
2026-09-24 01:24:11.123863 SELECT users.id AS users_id, users.name AS users_name, users.email AS users_email, users.passwo
2026-09-24 01:24:11.355619 SELECT users.id AS users_id, users.name AS users_name, users.email AS users_email, users.passwo
2026-09-24 01:24:11.470960 SELECT sessions.id AS sessions_id, sessions.user_id AS sessions_user_id, sessions.created_at AS
2026-09-24 01:24:11.474658 SELECT grocery_notices.id AS grocery_notices_id, grocery_notices.product_name AS grocery_notice
2026-09-24 01:24:11.478108 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.482907 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.486207 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.489019 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.492452 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.496261 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.499230 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.501568 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.503633 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.506730 SELECT notice_lots.id AS notice_lots_id, notice_lots.notice_id AS notice_lots_notice_id, notice
2026-09-24 01:24:11.590534 SELECT sessions.id AS sessions_id, sessions.user_id AS sessions_user_id, sessions.created_at AS
2026-09-24 01:24:11.595585 SELECT grocery_notices.id AS grocery_notices_id, grocery_notices.product_name AS grocery_notice
2026-09-24 01:24:11.599953 SELECT notice_lots.notice_id AS notice_lots_notice_id, notice_lots.id AS notice_lots_id, notice
$ mysql ... -e "SELECT LEFT(argument, 22) AS statement_start, COUNT(*) AS statements FROM mysql.general_log WHERE command_type = 'Query' AND
argument LIKE 'SELECT %' AND argument NOT LIKE '%general_log%' GROUP BY 1"
statement_start statements
SELECT users.id AS use  2
SELECT sessions.id AS  2
SELECT grocery_notices 2
SELECT notice_lots.id  10
SELECT notice_lots.not  1
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>

```

4, 5 and 6. Measurements

measure_n_plus_one.py logs in, sends 2 warm-up requests and then 30 measured requests for each of the six combinations, and stores every request (page size, version, request number, statement count, notices and lots returned, latency in ms, timestamp) in raw/n_plus_one_requests.csv (180 rows). The percentiles are computed with numpy.percentile and written to raw/n_plus_one_summary.json.

```
def timed_get(session, version, page_size):
    start = time.perf_counter()
    response = session.get(f"{BASE_URL}/api/notices/{version}", params={"page_size": page_size})
    latency_ms = (time.perf_counter() - start) * 1000
    response.raise_for_status()
    return response.json(), latency_ms
```

```
def main():
    session = requests.Session()
    session.post(f"{BASE_URL}/auth/login", json=LOGIN).raise_for_status()
```

```
    rows = []
    for page_size in PAGE_SIZES:
        for version in VERSIONS:
            for _ in range(WARMUP):
                timed_get(session, version, page_size)
            for number in range(1, REQUESTS_PER_CASE + 1):
                body, latency_ms = timed_get(session, version, page_size)
                rows.append({
                    "page_size": page_size,
                    "version": version,
                    "request": number,
                    "sql_queries": body["sql_queries"],
                    "notices_returned": len(body["notices"]),
                    "lots_returned": sum(len(notice["lots"]) for notice in body["notices"]),
                    "latency_ms": round(latency_ms, 3),
                    "timestamp": datetime.now().isoformat(timespec="milliseconds"),
                })
```

```
$ date
2026-09-23T18:22:18-07:00
$ python measure_n_plus_one.py
page_size=10 naive: 30 requests done, sql_queries=11
page_size=10 fixed: 30 requests done, sql_queries=2
page_size=50 naive: 30 requests done, sql_queries=51
page_size=50 fixed: 30 requests done, sql_queries=2
page_size=200 naive: 30 requests done, sql_queries=201
page_size=200 fixed: 30 requests done, sql_queries=2
```

	Page size	Version	SQL stmts/req	p50 (ms)	p95 (ms)	p99 (ms)
	10	naive	11	23.76	29.21	72.06
	10	fixed	2	10.12	12.04	12.81
	50	naive	51	92.52	149.12	160.95
	50	fixed	2	13.69	19.28	19.65
	200	naive	201	647.48	870.04	1050.93
	200	fixed	2	32.29	61.64	81.7

```
page_size=10: fixed is 2.35x faster at p50 (13.64 ms saved per request)
page_size=50: fixed is 6.76x faster at p50 (78.83 ms saved per request)
page_size=200: fixed is 20.05x faster at p50 (615.19 ms saved per request)
wrote ...reports\hw04\raw\n_plus_one_requests.csv (180 rows) and n_plus_one_summary.json
$ date
2026-09-23T18:22:45-07:00
```

```
cap-measure
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part
3: N+1 measurement)
=== Part 3: N+1 measurement (python measure_n_plus_one.py), server started with python app.py, Ollama idle ===
$ date
2026-09-23T18:22:18-07:00
$ python measure_n_plus_one.py
page_size=10 naive: 30 requests done, sql_queries=11
page_size=10 fixed: 30 requests done, sql_queries=2
page_size=50 naive: 30 requests done, sql_queries=51
page_size=50 fixed: 30 requests done, sql_queries=2
page_size=200 naive: 30 requests done, sql_queries=201
page_size=200 fixed: 30 requests done, sql_queries=2

| Page size | Version | SQL stmts/req | p50 (ms) | p95 (ms) | p99 (ms) |
|---|---|---|---|---|
| 10 | naive | 11 | 23.76 | 29.21 | 72.06 |
| 10 | fixed | 2 | 10.12 | 12.04 | 12.81 |
| 50 | naive | 51 | 92.52 | 149.12 | 160.95 |
| 50 | fixed | 2 | 13.69 | 19.28 | 19.65 |
| 200 | naive | 201 | 647.48 | 870.04 | 1050.93 |
| 200 | fixed | 2 | 32.29 | 61.64 | 81.7 |

page_size=10: fixed is 2.35x faster at p50 (13.64 ms saved per request)
page_size=50: fixed is 6.76x faster at p50 (78.83 ms saved per request)
page_size=200: fixed is 20.05x faster at p50 (615.19 ms saved per request)
wrote C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export\reports\hw04\raw\n_plus_one_requests.csv (180 rows)
and n_plus_one_summary.json
$ date
2026-09-23T18:22:45-07:00
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

Page size	Version	SQL stmts/req	p50 (ms)	p95 (ms)	p99 (ms)
10	naive	11	23.76	29.21	72.06
10	fixed	2	10.12	12.04	12.81
50	naive	51	92.52	149.12	160.95
50	fixed	2	13.69	19.28	19.65
200	naive	201	647.48	870.04	1050.93
200	fixed	2	32.29	61.64	81.70

7. How much faster the fixed version is, and why the speed-up grows

Page size	Statements removed	p50 naive / fixed	p95 naive / fixed	ms saved per request (p50)
10	9	2.35x	2.43x	13.64
50	49	6.76x	7.73x	78.83
200	199	20.05x	14.12x	615.19

The naive endpoint costs a fixed part plus one round trip per notice: the fixed part (session lookup, the notices query, building the response) is about 10 ms and each extra notice_lots statement adds roughly 1.5 ms at page sizes 10 and 50 and about 3 ms at 200, because every statement is a separate network round trip to MySQL (through the Docker port mapping here), a separate parse and execute on the server, and a separate set of ORM objects. The fixed endpoint always runs two statements, and the second one returns all lots of the page at once, so its cost grows only with the rows it returns (10 ms at page size 10, 32 ms at 200). The saving is therefore proportional to the page size while the denominator hardly moves: at 10 rows the fixed part still dominates both versions and the ratio is only

2.4x; at 200 rows the naive version spends almost all of its 650 ms in the 200 extra round trips and the ratio reaches 20x. The p99 of the naive version at page size 10 (72 ms) is one slow request among the 30 (89 ms in the raw file), which is why the p95 column is the fairer comparison of the tails.

8. One index, EXPLAIN before and after

The index is on `grocery_notices.notice_source` (migrations/hw04_add_index.sql), the column the seeded sources are filtered by; the query is `SELECT id, product_name, notice_source FROM grocery_notices WHERE notice_source = 'FDA recall bulletin' ORDER BY id`.

```
CREATE INDEX ix_grocery_notices_notice_source ON grocery_notices (notice_source);
```

Before:

```
$ mysql ... -e "EXPLAIN SELECT id, product_name, notice_source FROM grocery_notices WHERE
notice_source = 'FDA recall bulletin' ORDER BY id"
id select_type table partitions type possible_keys key key_len ref rows
filtered Extra
1 SIMPLE grocery_notices NULL index NULL PRIMARY 4 NULL 5158
10.00 Using where

$ mysql ... -e "EXPLAIN ANALYZE ..."
-> Filter: (grocery_notices.notice_source = 'FDA recall bulletin') (cost=522 rows=516) (actual
time=0.0767..2.16 rows=183 loops=1)
-> Index scan on grocery_notices using PRIMARY (cost=522 rows=5158) (actual time=0.0396..1.69
rows=5000 loops=1)
```

```
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "DROP INDEX ix_grocery_notices_notice_source ON grocery_notices" (table as it was before step 8)
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "SELECT INDEX_NAME, COLUMN_NAME, INDEX_TYPE, CARDINALITY FROM information_schema.STATISTICS WHERE TABLE_SCHEMA = 's0571_rel' AND TABLE_NAME = 'grocery_notices'"
+-----+-----+-----+-----+
| INDEX_NAME | COLUMN_NAME | INDEX_TYPE | CARDINALITY |
+-----+-----+-----+-----+
| PRIMARY    | id          | BTREE     | 5158        |
+-----+-----+-----+-----+
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "EXPLAIN SELECT id, product_name, notice_source FROM grocery_notices WHERE notice_source = 'FDA recall bulletin' ORDER BY id"
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table      | partitions | type | possible_keys | key      | key_len | ref | rows | filtered | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1  | SIMPLE      | grocery_notices | NULL       | index | NULL          | PRIMARY | 4       | NULL | 5158 | 10.00    | Using where |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "EXPLAIN ANALYZE SELECT id, product_name, notice_source FROM grocery_notices WHERE notice_source = 'FDA recall bulletin' ORDER BY id"
+-----+-----+-----+-----+
| EXPLAIN |
+-----+-----+-----+-----+
|         |
+-----+-----+-----+-----+
| -> Filter: (grocery_notices.notice_source = 'FDA recall bulletin') (cost=522 rows=516) (actual time=0.101..1.47 rows=183 loops=1)
| -> Index scan on grocery_notices using PRIMARY (cost=522 rows=5158) (actual time=0.0458..1.13 rows=5000 loops=1)
|         |
+-----+-----+-----+-----+
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

After:

```
$ mysql ... s0571_rel < migrations/hw04_add_index.sql
$ mysql ... -e "EXPLAIN SELECT id, product_name, notice_source FROM grocery_notices WHERE
notice_source = 'FDA recall bulletin' ORDER BY id"
id select_type table partitions type possible_keys key
key_len ref rows filtered Extra
1 SIMPLE grocery_notices NULL ref ix_grocery_notices_notice_source
ix_grocery_notices_notice_source 802 const 183 100.00 NULL
$ mysql ... -e "EXPLAIN ANALYZE ..."
```

-> Index lookup on grocery_notices using ix_grocery_notices_notice_source (notice_source='FDA recall bulletin') (cost=37 rows=183) (actual time=0.0887..0.593 rows=183 loops=1)

```
cap-explain-after
PS> Get-Content migrations\hw04_add_index.sql
CREATE INDEX ix_grocery_notices_notice_source ON grocery_notices (notice_source);
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "source migrations/hw04_add_index.sql"
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "SELECT INDEX_NAME, COLUMN_NAME, INDEX_TYPE, CARDINALITY FROM information_schema.STATISTICS WHERE TABLE_SCHEMA = 's0571_rel' AND TABLE_NAME = 'grocery_notices'"
+-----+-----+-----+-----+
| INDEX_NAME | COLUMN_NAME | INDEX_TYPE | CARDINALITY |
+-----+-----+-----+-----+
| ix_grocery_notices_notice_source | notice_source | BTREE | 24 |
| PRIMARY | id | BTREE | 5158 |
+-----+-----+-----+-----+
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "EXPLAIN SELECT id, product_name, notice_source FROM grocery_notices WHERE notice_source = 'FDA recall bulletin' ORDER BY id"
+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | row |
|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 1 | SIMPLE | grocery_notices | NULL | ref | ix_grocery_notices_notice_source | ix_grocery_notices_notice_source | 802 | const | 183 |
| 100.00 | NULL |
+-----+-----+-----+-----+
PS> mysql -h 127.0.0.1 -P 3307 -u root -p s0571_rel -e "EXPLAIN ANALYZE SELECT id, product_name, notice_source FROM grocery_notices WHERE notice_source = 'FDA recall bulletin' ORDER BY id"
+-----+
| EXPLAIN |
+-----+
| -> Index lookup on grocery_notices using ix_grocery_notices_notice_source (notice_source='FDA recall bulletin') (cost=37 rows=183) (actual time=0.0887..0.593 rows=183 loops=1) |
+-----+
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_ANS_Docker_export>
```

What changed: before the index there was no possible_keys entry for the WHERE column, so MySQL scanned the whole table (type=index over PRIMARY, 5,158 estimated rows, 5,000 actually read) and applied the filter to every row (Using where, filtered=10.00), keeping 183. After the index the access type is ref with ref=const: the optimizer looks up 'FDA recall bulletin' in the B-tree and reads only the 183 matching entries, the row estimate equals the real count, filtered is 100 percent, and the Using where step is gone. EXPLAIN ANALYZE shows the cost estimate dropping from 522 to 37 and the actual time from 2.16 ms to 0.59 ms. key_len=802 is the 200-character utf8mb4 column (4 bytes per character plus 2 length bytes). The ORDER BY id needs no filesort in either plan: the primary-key scan is already in id order, and an InnoDB secondary index stores the primary key with every entry, so the entries for one source also come out in id order.

9. Both versions of the endpoint in the API client at each page size

Each response shows version, page_size, sql_queries and the notices with their lots.

Page size	Naive	Fixed
10	api-naive-10.png (sql_queries: 11)	api-fixed-10.png (sql_queries: 2)
50	api-naive-50.png (sql_queries: 51)	api-fixed-50.png (sql_queries: 2)
200	api-naive-200.png (sql_queries: 201)	api-fixed-200.png (sql_queries: 2)

GET

/api/notices/naive

List Naive

⌵

Parameters

Cancel

Name	Description
page_size	
integer	10
(query)	

Execute

Clear

Responses

Curl

curl -X 'GET' \n'http://localhost:8571/api/notices/naive?page_size=10' \n-H 'accept: */*' \n

Request URL

http://localhost:8571/api/notices/naive?page_size=10

Server response

Code	Details
200	Response body <pre>{ "version": "naive", "page_size": 10, "sql_queries": 11, "notices": [{ "productName": "Dried tahini 1 gallon", "noticeSource": "Abbott Nutrition", "id": 5000, "lots": [{ "id": 183, "lotCode": "LOT-553350", "unitsAffected": 2298 }] }, { "productName": "Fresh chicken thighs 12 oz", "noticeSource": "Harvest Fresh Produce", "id": 4999, "lots": [{ "id": 4, "lotCode": "LOT-951024", "unitsAffected": 1585 }] }] }</pre> Download

Response headers

content-length: 1527\ncontent-type: application/json

GET

/api/notices/fixed

List Fixed

^

Parameters

Cancel

Name	Description
page_size	integer (query)
10	

Execute

Clear

Responses

Curl

curl -X 'GET' \ 'http://localhost:8571/api/notices/fixed?page_size=10' \ -H 'accept: */*' 

Request URL

http://localhost:8571/api/notices/fixed?page_size=10

Server response

Code	Details
200	Response body <pre>{ "version": "fixed", "page_size": 10, "sql_queries": 2, "notices": [{ "productName": "Dried tahini 1 gallon", "noticeSource": "Abbott Nutrition", "id": 5000, "lots": [{ "id": 183, "lotCode": "LOT-553350", "unitsAffected": 2298 }] }, { "productName": "Fresh chicken thighs 12 oz", "noticeSource": "Harvest Fresh Produce", "id": 4999, "lots": [{ "id": 4, "lotCode": "LOT-951024", "unitsAffected": 1585 }] }] }</pre>  Download Response headers content-length: 1526 content-type: application/json

GET

/api/notices/naive

List Naive

^

Parameters

Cancel

Name	Description
page_size	
integer	50
(query)	

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
'http://localhost:8571/api/notices/naive?page_size=50' \
-H 'accept: */*'
```

Request URL

```
http://localhost:8571/api/notices/naive?page_size=50
```

Server response

Code	Details
200	Response body <pre>{ "version": "naive", "page_size": 50, "sql_queries": 51, "notices": [{ "productName": "Dried tahini 1 gallon", "noticeSource": "Abbott Nutrition", "id": 5000, "lots": [{ "id": 183, "lotCode": "LOT-553350", "unitsAffected": 2298 }] }, { "productName": "Fresh chicken thighs 12 oz", "noticeSource": "Harvest Fresh Produce", "id": 4999, "lots": [{ "id": 4, "lotCode": "LOT-951024", "unitsAffected": 1585 }] }] }</pre> Download

Response headers

```
content-length: 8087
content-type: application/json
```

GET

/api/notices/fixed

List Fixed

^

Parameters

Cancel

Name	Description
page_size	integer (query)
	50

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
'http://localhost:8571/api/notices/fixed?page_size=50' \
-H 'accept: */*'
```

Request URL

```
http://localhost:8571/api/notices/fixed?page_size=50
```

Server response

Code	Details
200	Response body <pre>{ "version": "fixed", "page_size": 50, "sql_queries": 2, "notices": [{ "productName": "Dried tahini 1 gallon", "noticeSource": "Abbott Nutrition", "id": 5000, "lots": [{ "id": 183, "lotCode": "LOT-553350", "unitsAffected": 2298 }] }, { "productName": "Fresh chicken thighs 12 oz", "noticeSource": "Harvest Fresh Produce", "id": 4999, "lots": [{ "id": 4, "lotCode": "LOT-951024", "unitsAffected": 1585 }] }] }</pre> Download

Response headers

```
content-length: 8086
content-type: application/json
```

GET

/api/notices/naive

List Naive

^

Parameters

Cancel

Name	Description
page_size	
integer	200
(query)	

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
'http://localhost:8571/api/notices/naive?page_size=200' \
-H 'accept: */*'
```

Request URL

```
http://localhost:8571/api/notices/naive?page_size=200
```

Server response

Code	Details
200	Response body <pre>{ "version": "naive", "page_size": 200, "sql_queries": 201, "notices": [{ "productName": "Dried tahini 1 gallon", "noticeSource": "Abbott Nutrition", "id": 5000, "lots": [{ "id": 183, "lotCode": "LOT-553350", "unitsAffected": 2298 }] }, { "productName": "Fresh chicken thighs 12 oz", "noticeSource": "Harvest Fresh Produce", "id": 4999, "lots": [{ "id": 4, "lotCode": "LOT-951024", "unitsAffected": 1585 }] }] }</pre> Download

Response headers

```
content-length: 30976
content-type: application/json
```

GET
/api/notices/fixed
List Fixed

Parameters
Cancel

Name	Description
page_size integer (query)	200

Execute
Clear

Responses

Curl

```
curl -X 'GET' \
'http://localhost:8571/api/notices/fixed?page_size=200' \
-H 'accept: */*'

```

Request URL

```
http://localhost:8571/api/notices/fixed?page_size=200

```

Server response

Code	Details
200	Response body <pre>{ "version": "fixed", "page_size": 200, "sql_queries": 2, "notices": [{ "productName": "Dried tahini 1 gallon", "noticeSource": "Abbott Nutrition", "id": 5000, "lots": [{ "id": 183, "lotCode": "LOT-553350", "unitsAffected": 2298 }] }, { "productName": "Fresh chicken thighs 12 oz", "noticeSource": "Harvest Fresh Produce", "id": 4999, "lots": [{ "id": 4, "lotCode": "LOT-951024", "unitsAffected": 1585 }] }] }</pre> Download Response headers <pre>content-length: 30974 content-type: application/json </pre>

The same two responses from curl at page size 10 (RUN_LOG.txt):

```
$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices/naive?page_size=10" | python -c "...
{'version': 'naive', 'page_size': 10, 'sql_queries': 11} first notice: {'productName': 'Dried
tahini 1 gallon', 'noticeSource': 'Abbott Nutrition', 'id': 5000, 'lots': [{'id': 183, 'lotCode':
'LOT-553350', 'unitsAffected': 2298}]}"
$ curl -s -b cookies.txt "http://127.0.0.1:8571/api/notices/fixed?page_size=10" | python -c "...
{'version': 'fixed', 'page_size': 10, 'sql_queries': 2} first notice: {'productName': 'Dried tahini
1 gallon', 'noticeSource': 'Abbott Nutrition', 'id': 5000, 'lots': [{'id': 183, 'lotCode': 'LOT-
553350', 'unitsAffected': 2298}]}"
```

Part 4: grounded RAG question answering

Everything is in rag.py (python rag.py, about 13 minutes on this laptop; python rag.py --summary re-evaluates the saved answers and rebuilds the tables). The model is qwen3:8b through

Ollama with temperature 0, thinking off, num_ctx 4096 and at most 200 output tokens, called through src/model_client.py from HW1. The retrieved-chunk printouts, the answers, the comparison, the k sweep and the evaluation are saved in reports/hw04/raw/ (rag_retrievals.txt, rag_results.jsonl, rag_comparison.md, rag_k_sweep.md, rag_evaluation.md, plus rag_chunks.jsonl with every chunk).


1. Corpus and index

The corpus is the 21 plain-text documents in corpus/ collected for HW3 (FDA recall pages, 21 CFR Part 7 Subpart C, the FSMA and Food Traceability Rule pages, openFDA enforcement reports, a USDA page and 13 Wikipedia articles on recalls, outbreaks and food distribution; 393,964 bytes). Each file is split with RecursiveCharacterTextSplitter(chunk_size=500, chunk_overlap=50); every chunk keeps its text, source file name and a running chunk_id. The chunks are embedded with all-MiniLM-L6-v2 (normalised, so the inner product is the cosine similarity) and loaded into a FAISS IndexFlatIP.

```
def load_chunks():
    splitter = RecursiveCharacterTextSplitter(chunk_size=CHUNK_SIZE, chunk_overlap=CHUNK_OVERLAP)
    chunks = []
    for path in sorted(CORPUS.glob("*.txt")):
        for text in splitter.split_text(path.read_text(encoding="utf-8")):
            chunks.append({"chunk_id": len(chunks), "source": path.name, "text": text})
    return chunks

def build_index(chunks, embedder):
    vectors = embedder.encode([chunk["text"] for chunk in chunks], normalize_embeddings=True,
batch_size=64)
    index = faiss.IndexFlatIP(vectors.shape[1])
    index.add(np.asarray(vectors, dtype="float32"))
    return index

$ python rag.py
21 documents, 1158 chunks (chunk_size=500, chunk_overlap=50)
FAISS index built with 1158 vectors of dimension 384
```



```
cap-rag-index
PS C:\Users\Poushall\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part 4: RAG run)
=== Part 4: RAG run (PYTHONUTF8=1 python rag.py), Ollama qwen3:8b, started 2026-09-23T18:24:33-07:00, last answer written 2026-09-23T18:37:44-07:00 ===
$ python rag.py
21 documents, 1158 chunks (chunk_size=500, chunk_overlap=50)
FAISS index built with 1158 vectors of dimension 384
[tokens] turn=1 input=58 output=53 total=111
[Q1 A k=3] The 2008 listeriosis outbreak in Canada was linked to the Maple Leaf Foods facility in Toronto, Ontario. This outbreak resulted in multiple deaths and led to a widespread recall of products. It was the largest listeriosis outbreak in Canadian history.

--- Q1 config B k=3: top-3 retrieved chunks ---
1. score=0.8481 source=wikipedia_2008_canadian_listeriosis_outbreak.txt chunk_id=501 The 2008 Canadian listeriosis outbreak was a widespread outbreak of listeriosis in Canada linked to cold cuts from a Maple Leaf Foods plant
2. score=0.8144 source=wikipedia_2008_canadian_listeriosis_outbreak.txt chunk_id=502 Listeriosis is an infection caused by the bacterium Listeria monocytogenes. The outbreak originated from lines 8 and 9 of the Maple Leaf Foods
3. score=0.6580 source=wikipedia_2008_canadian_listeriosis_outbreak.txt chunk_id=507 Officials from Maple Leaf believe that the outbreak originated sometime in July 2008 on line 8 or line 9 of the North York facility. Regarding
[tokens] turn=2 input=337 output=80 total=417
PS C:\Users\Poushall\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

2. Retrieval with printed chunks

retrieve returns the top-k chunks with their score, and show_hits prints rank, score, source, chunk id and the first 140 characters of every chunk before the model is called (and appends the same lines to raw/rag_retrievals.txt), so a bad retrieval can be told apart from a bad generation.

```
def retrieve(question, k, chunks, index, embedder):
    vector = embedder.encode([question], normalize_embeddings=True)
```

```

scores, ids = index.search(np.asarray(vector, dtype="float32"), k)
return [dict(chunks[i], score=round(float(score), 4)) for score, i in zip(scores[0], ids[0])]

def show_hits(label, hits):
    lines = [f"--- {label}: top-{len(hits)} retrieved chunks ---"]
    for rank, hit in enumerate(hits, 1):
        preview = " ".join(hit["text"].split())[:140]
        lines.append(f" {rank}. score={hit['score']:.4f} source={hit['source']} chunk_id={hit['chunk_id']} {preview}")
    text = "\n".join(lines)
    print(text)
    with (RAW / "rag_retrievals.txt").open("a", encoding="utf-8") as f:
        f.write(text + "\n\n")

```

Printouts for Q1 (right chunks), Q2 (wrong chunks: all from one document, two of them heading fragments) and Q6 (unrelated question, low scores):

```

--- Q1 config B k=3: top-3 retrieved chunks ---
1. score=0.8481 source=wikipedia_2008_canadian_listeriosis_outbreak.txt chunk_id=501 The 2008
Canadian listeriosis outbreak was a widespread outbreak of listeriosis in Canada linked to cold
cuts from a Maple Leaf Foods plant
2. score=0.8144 source=wikipedia_2008_canadian_listeriosis_outbreak.txt chunk_id=502
Listeriosis is an infection caused by the bacterium Listeria monocytogenes . The outbreak
originated from lines 8 and 9 of the Maple Leaf Fo
3. score=0.6580 source=wikipedia_2008_canadian_listeriosis_outbreak.txt chunk_id=507 Officials
from Maple Leaf believe that the outbreak originated sometime in July 2008 on line 8 or line 9 of
the North York facility. Regardl

--- Q2 config B k=3: top-3 retrieved chunks ---
1. score=0.7282 source=fda_fsma_food_traceability_rule.txt chunk_id=153 in § 1.1315 of the
Food Traceability Rule and reflect the current practices specific to the covered entity.
2. score=0.7088 source=fda_fsma_food_traceability_rule.txt chunk_id=95 Getting Started with
the Food Traceability Rule
3. score=0.7071 source=fda_fsma_food_traceability_rule.txt chunk_id=94 proposed to extend the
compliance date for the rule by 30 months to July 20, 2028. Subsequently, the Continuing
Appropriations, Agriculture,

--- Q6 config C k=3: top-3 retrieved chunks ---
1. score=0.3820 source=wikipedia_food_distribution.txt chunk_id=817 Japan
2. score=0.2249 source=wikipedia_food_distribution.txt chunk_id=796 Farm to Table
3. score=0.2023 source=wikipedia_food_distribution.txt chunk_id=816 See also
dropped chunk 796: score 0.2249 below 0.35
dropped chunk 816: score 0.2023 below 0.35
context after filtering: [817]

```

```
cap-rag-retrievals
outbreak originated from lines 8 and 9 of the Maple Leaf Fo
3. score=0.6588 source=wikipedia_2008_canadian_listeria_outbreak.txt chunk_id=507 Officials from Maple Leaf believe that the outbreak originated sometime in July
2008 on line 8 or line 9 of the North York facility. Regardl

--- Q1 config C k=3: top-3 retrieved chunks ---
1. score=0.8481 source=wikipedia_2008_canadian_listeria_outbreak.txt chunk_id=501 The 2008 Canadian listeria outbreak was a widespread outbreak of listeria
in Canada linked to cold cuts from a Maple Leaf Foods plant
2. score=0.8144 source=wikipedia_2008_canadian_listeria_outbreak.txt chunk_id=502 Listeria is an infection caused by the bacterium Listeria monocytogenes . The
outbreak originated from lines 8 and 9 of the Maple Leaf Fo
3. score=0.6588 source=wikipedia_2008_canadian_listeria_outbreak.txt chunk_id=507 Officials from Maple Leaf believe that the outbreak originated sometime in July
2008 on line 8 or line 9 of the North York facility. Regardl

--- Q2 config B k=3: top-3 retrieved chunks ---
1. score=0.7282 source=fda_fsma_food_traceability_rule.txt chunk_id=153 in § 1.1315 of the Food Traceability Rule and reflect the current practices specific to the c
overed entity.
2. score=0.7888 source=fda_fsma_food_traceability_rule.txt chunk_id=95 Getting Started with the Food Traceability Rule
3. score=0.7871 source=fda_fsma_food_traceability_rule.txt chunk_id=94 proposed to extend the compliance date for the rule by 30 months to July 20, 2028. Subsequentl
y, the Continuing Appropriations, Agriculture,

--- Q2 config C k=3: top-3 retrieved chunks ---
1. score=0.7282 source=fda_fsma_food_traceability_rule.txt chunk_id=153 in § 1.1315 of the Food Traceability Rule and reflect the current practices specific to the c
overed entity.
2. score=0.7888 source=fda_fsma_food_traceability_rule.txt chunk_id=95 Getting Started with the Food Traceability Rule
3. score=0.7871 source=fda_fsma_food_traceability_rule.txt chunk_id=94 proposed to extend the compliance date for the rule by 30 months to July 20, 2028. Subsequentl
y, the Continuing Appropriations, Agriculture,

--- Q3 config B k=3: top-3 retrieved chunks ---
1. score=0.6841 source=wikipedia_peanut_corporation_of_america.txt chunk_id=979 "costing us huge $$$$." In a June 2008 email exchange, Parnell complained to a worke
r after being notified that Salmonella had been found i
2. score=0.6499 source=wikipedia_peanut_corporation_of_america.txt chunk_id=972 analysis and laboratory testing by state officials in Minnesota and Connecticut , the
Food and Drug Administration (FDA) and the CDC confirm
3. score=0.6366 source=wikipedia_peanut_corporation_of_america.txt chunk_id=964 In late 2008 and early 2009, as a result of the Salmonella contamination event, nine
people died and at least 714 people fell ill from food

--- Q3 config C k=3: top-3 retrieved chunks ---
1. score=0.6841 source=wikipedia_peanut_corporation_of_america.txt chunk_id=979 "costing us huge $$$$." In a June 2008 email exchange, Parnell complained to a worke
r after being notified that Salmonella had been found i
2. score=0.6499 source=wikipedia_peanut_corporation_of_america.txt chunk_id=972 analysis and laboratory testing by state officials in Minnesota and Connecticut , the
Food and Drug Administration (FDA) and the CDC confirm
3. score=0.6366 source=wikipedia_peanut_corporation_of_america.txt chunk_id=964 In late 2008 and early 2009, as a result of the Salmonella contamination event, nine
people died and at least 714 people fell ill from food
```

3. Three configurations

(A) sends the question alone. (B) puts the three raw chunks in front of the question. (C) drops chunks with a cosine score below 0.35 and chunks whose word set overlaps a kept chunk by more than 60 percent, keeps the survivors in score order, labels each one [Source n] file, chunk id, and adds the grounding rules: answer only from the sources, cite the source number, and reply with the exact sentence "I cannot answer this question from the provided documents" when the sources are not enough.

```
NO_RAG_PROMPT = "You are a helpful assistant. Answer the question in at most three sentences."
BASIC_PROMPT = "Use the context to answer the question in at most three sentences."
CONTEXT_PROMPT = (
    "Answer the question using only the numbered sources below, in at most three sentences."
    "Cite the source number after each fact, like [Source 2]. "
    f"If the sources do not contain the answer, reply exactly: {REFUSAL}"
)

def engineer_context(hits):
    kept, dropped = [], []
    for hit in hits:
        if hit["score"] < MIN_SCORE:
            dropped.append({"chunk_id": hit["chunk_id"], "reason": f"score {hit['score']:.4f} below {MIN_SCORE}"})
        elif any(overlap(hit["text"], other["text"]) > DUPLICATE_OVERLAP for other in kept):
            dropped.append({"chunk_id": hit["chunk_id"], "reason": "duplicate of a kept chunk"})
        else:
            kept.append(hit)
    return kept, dropped

def basic_prompt(question, hits):
    context = "\n\n".join(hit["text"] for hit in hits)
    return f"Context:\n{context}\n\nQuestion: {question}"

def engineered_prompt(question, kept):
```

```

parts = [f"[Source {i}] {hit['source']}, chunk {hit['chunk_id']}\n{hit['text']}" for i, hit in
enumerate(kept, 1)]
return "Sources:\n\n" + "\n\n".join(parts) + f"\n\nQuestion: {question}"

```

4. The six questions

Id	Kind	Question	Where the answer is
Q1	answer in one chunk	Which Maple Leaf Foods facility was the source of the 2008 listeriosis outbreak in Canada?	wikipedia_2008_canadian_listeriosis_outbreak.txt, chunk 502 ("Bartor Road facility, Establishment No. 97B")
Q2	answer needs two chunks	When was the Food Safety Modernization Act signed into law, and what was the original compliance date of the Food Traceability Rule?	wikipedia_fda_food_safety_modernization_act.txt chunk 673 (January 4, 2011) and fda_fsma_food_traceability_rule.txt chunk 93 (January 20, 2026)
Q3	similar information across documents	Which company was responsible for the 2008 to 2009 Salmonella outbreak linked to peanut butter, and what happened to the company afterwards?	wikipedia_peanut_corporation_of_america.txt (chunks 964, 965, 972) and wikipedia_product_recall.txt (chunk 1047) both describe PCA, the recall and the shutdown
Q4	ambiguous	What caused the outbreak?	Five outbreaks in the corpus (spinach 2006, Canada 2008, PCA 2008 to 2009, Germany 2011, formula 2022); the good answer asks which one
Q5	answer not in the documents	What is the FDA's toll-free phone number for reporting a problem with a recalled food?	Not in the corpus (no phone number appears in any file); must be refused
Q6	unrelated	Who won the 2018 FIFA World Cup?	Not in the corpus; must be refused

Answers of the three configurations (full text in `raw/rag_comparison.md`; "n/a" means not applicable):

Q	(A) No RAG	(B) Basic RAG, top-3 raw chunks	(C) Context-engineered RAG
Q1	"linked to the Maple Leaf Foods facility in Toronto, Ontario ... the largest listeriosis outbreak in Canadian history" (no plant named, wrong)	"originated from lines 8 and 9 of the Maple Leaf Foods Bortor Road facility in Toronto, Ontario ... establishment number 97B ... 57 confirmed cases and 23 deaths" (correct)	"the Maple Leaf Foods plant in Toronto, Ontario [Source 1]. Specifically ... lines 8 and 9 of the Maple Leaf Foods Bortor Road facility (Establishment No. 97B) [Source 2]" (correct, cited)
Q2	"signed into law on January 4, 2011 ... original compliance date ... January 1, 2022 ... extended to January 1, 2025" (second half invented)	"signed into law in 2011 ... original compliance date ... July 20, 2026 ... extended to July 20, 2028" (date wrong, not in the chunks)	"I cannot answer this question from the provided documents." (evidence really was missing)
Q3	"Peanut Corporation of America (PCA) ... massive recall and eventual bankruptcy ... Stewart Parnell was convicted" (correct, from memory)	"Peanut Corporation of America (PCA) ... most extensive food recall in U.S. history ... nine deaths and over 700 illnesses ... the company's bankruptcy" (correct, but "bankruptcy" is not in the three chunks)	"Peanut Corporation of America (PCA) ... [Source 2]. ... most extensive food recall in U.S. history ... 46 states, more than 360 companies, more than 3,900 different products [Source 3]" (grounded but the second half of the question is not answered)
Q4	"depends on the specific disease or event in question ... more context is needed about the outbreak being referred to" (asks for clarification, the right reaction)	"likely caused by contaminated spinach ... Most victims became infected in Germany or France" (mixes the 2006 spinach and the 2011 German outbreak into one)	"contaminated food, specifically spinach ... [Source 2] ... infected in Germany or France [Source 3]. The exact cause was not specified in the provided documents." (same mix, hedged, not the exact refusal)
Q5	"1-800-FDA-1040 (1-800-323-1040) ... FDA's Food Emergency Response Team" (invented)	"The context provided does not include the FDA's toll-free phone number ... consumers should visit the FDA's website" (soft refusal)	"I cannot answer this question from the provided documents."
Q6	"France won the 2018 FIFA World Cup, defeating Croatia 4-2 ... after a penalty shootout in extra time" (answered from memory, with an invented detail)	"The 2018 FIFA World Cup was won by France ... defeated Croatia in the final" (ignored the context, answered from memory)	"I cannot answer this question from the provided documents."

Q5 and Q6 are refused by configuration C with the exact sentence; configuration B refused Q5 softly but answered Q6, and configuration A answered both.

```
cap-rag-answers
[tokens] turn=1 input=58 output=53 total=111
[Q1 A k=3] The 2008 listeriosis outbreak in Canada was linked to the Maple Leaf Foods facility in Toronto, Ontario. This outbreak resulted in multiple deaths and a widespread recall of products. It was the largest listeriosis outbreak in Canadian history.

--- Q1 config B k=3: top-3 retrieved chunks ---
1. score=0.8481 source=Wikipedia:2008_canadian_listeriosis_outbreak.txt chunk_id=501 The 2008 Canadian listeriosis outbreak was a widespread outbreak in Canada linked to cold cuts from a Maple Leaf Foods plant
2. score=0.8144 source=Wikipedia:2008_canadian_listeriosis_outbreak.txt chunk_id=502 Listeriosis is an infection caused by the bacterium Listeria monocytogenes. The outbreak originated from lines 8 and 9 of the Maple Leaf Foods facility in Toronto, Ontario.
3. score=0.6580 source=Wikipedia:2008_canadian_listeriosis_outbreak.txt chunk_id=507 Officials from Maple Leaf believe that the outbreak originated from line 8 or line 9 of the North York facility. Regarding the contamination of cold cuts, the company stated that the bacteria likely contaminated the packaging process.

[tokens] turn=2 input=337 output=88 total=417
[Q1 B k=3] The 2008 listeriosis outbreak in Canada originated from lines 8 and 9 of the Maple Leaf Foods Bantrock Road facility in Toronto, Ontario. This facility was established number 978, was linked to the contamination of cold cuts that led to 57 confirmed cases and 23 deaths. The bacteria likely contaminated the packaging process.

--- Q1 config C k=3: top-3 retrieved chunks ---
1. score=0.8481 source=Wikipedia:2008_canadian_listeriosis_outbreak.txt chunk_id=501 The 2008 Canadian listeriosis outbreak was a widespread outbreak in Canada linked to cold cuts from a Maple Leaf Foods plant
2. score=0.8144 source=Wikipedia:2008_canadian_listeriosis_outbreak.txt chunk_id=502 Listeriosis is an infection caused by the bacterium Listeria monocytogenes. The outbreak originated from lines 8 and 9 of the Maple Leaf Foods facility in Toronto, Ontario.
3. score=0.6580 source=Wikipedia:2008_canadian_listeriosis_outbreak.txt chunk_id=507 Officials from Maple Leaf believe that the outbreak originated from line 8 or line 9 of the North York facility. Regarding the contamination of cold cuts, the company stated that the bacteria likely contaminated the packaging process.

[tokens] turn=3 input=454 output=69 total=523
[Q1 C k=3] The source of the 2008 listeriosis outbreak was the Maple Leaf Foods plant in Toronto, Ontario [Source 1]. Specifically, the outbreak originated from lines 8 and 9 of the Maple Leaf Foods Bantrock Road facility (Establishment No. 978) in Toronto, Ontario [Source 2].

[tokens] turn=4 input=62 output=79 total=141
[Q2 A k=3] The Food Safety Modernization Act (FSMA) was signed into law on January 4, 2011. The original compliance date for the Food Traceability Rule, FSMA, was January 1, 2022. However, the compliance date was later extended to January 1, 2025, for most covered products.

--- Q2 config B k=3: top-3 retrieved chunks ---
1. score=0.7381 source=FDA:Food_Safety_Modernization_Act_FSMMA.txt chunk_id=152 In 2011, the Food Safety Modernization Act (FSMA) was signed into law, and it reflected the current practice of food safety. The act was designed to prevent foodborne illness and protect public health. The act was signed into law on January 4, 2011. The original compliance date for the Food Traceability Rule, FSMA, was January 1, 2022. However, the compliance date was later extended to January 1, 2025, for most covered products.
```

5. Context size sweep (Q2 at $k = 1, 3, 5$, configuration C)

k	Retrieved (source, chunk, score)	Kept	Answer	Correct	Grounded
1	traceability #153 (0.7282)	153	"I cannot answer this question from the provided documents."	no	n/a
3	traceability #153 (0.7282), #95 (0.7088), #94 (0.7071)	153, 95, 94	"signed into law on January 4, 2011 [Source 3]. The original compliance date ... was July 20, 2026, but it was extended to July 20, 2028 [Source 3]."	no	no (Source 3 contains neither date; the compliance date is wrong)
5	traceability #153, #95, #94, #93 (0.7028), #169 (0.6987)	153, 95, 94, 93, 169	"signed into law on January 4, 2011 [Source 3]. The original compliance date of the Food Traceability Rule was January 20, 2026 [Source 4]."	yes	no (January 20, 2026 is in Source 4, but January 4, 2011 is not in Source 3)

More context helped: k=1 gave one fragment (107 characters about § 1.1315) and the model refused; k=5 finally brought chunk 93 with the compliance date and the answer became correct. Irrelevant chunks appeared at every k: #153 and #95 are a sentence fragment and a heading, and #169 (rank 5) is about a public meeting. The FSMA signing date was never retrieved at any k, because the question embeds close to "Food Traceability Rule" and every hit comes from that page; the model filled that half in from memory and attached a citation to a chunk that does not contain it. k=5 was the best of the three for accuracy, but the extra context did not make the answer more faithful, and the same k=3 prompt produced a refusal in the comparison run and an answer here (temperature 0, identical prompt of 283 tokens), so the k=3 row is also a reminder that the model is not deterministic.

```
cap-rag-sweep
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Part 4: RAG run)

--- Q2 config C k=1: top-1 retrieved chunks ---
1. score=0.7282 source=fda_fsma_food_traceability_rule.txt chunk_id=153 in § 1.1315 of the Food Traceability Rule and reflect the current practices specific to the covered entity.
context after filtering: [153]
[tokens] turn=19 input=150 output=11 total=161
[Q2 C k=1] I cannot answer this question from the provided documents.

--- Q2 config C k=3: top-3 retrieved chunks ---
1. score=0.7282 source=fda_fsma_food_traceability_rule.txt chunk_id=153 in § 1.1315 of the Food Traceability Rule and reflect the current practices specific to the covered entity.
2. score=0.7088 source=fda_fsma_food_traceability_rule.txt chunk_id=95 Getting Started with the Food Traceability Rule
3. score=0.7071 source=fda_fsma_food_traceability_rule.txt chunk_id=94 proposed to extend the compliance date for the rule by 30 months to July 20, 2028. Subsequently, the Continuing Appropriations, Agriculture,
context after filtering: [153, 95, 94]
[tokens] turn=20 input=283 output=68 total=351
[Q2 C k=3] The Food Safety Modernization Act was signed into law on January 4, 2011 [Source 3]. The original compliance date of the Food Traceability Rule was July 20, 2026, but it was extended to July 20, 2028 [Source 3].

--- Q2 config C k=5: top-5 retrieved chunks ---
1. score=0.7282 source=fda_fsma_food_traceability_rule.txt chunk_id=153 in § 1.1315 of the Food Traceability Rule and reflect the current practices specific to the covered entity.
2. score=0.7088 source=fda_fsma_food_traceability_rule.txt chunk_id=95 Getting Started with the Food Traceability Rule
3. score=0.7071 source=fda_fsma_food_traceability_rule.txt chunk_id=94 proposed to extend the compliance date for the rule by 30 months to July 20, 2028. Subsequently, the Continuing Appropriations, Agriculture,
4. score=0.7028 source=fda_fsma_food_traceability_rule.txt chunk_id=93 Compliance Date Because the Food Traceability Final Rule requires entities to share information with other entities in their supply chain, the
5. score=0.6987 source=fda_fsma_food_traceability_rule.txt chunk_id=169 The Reagan-Udall Foundation held a virtual public meeting to hear insights on the Food Traceability Rule from both FDA leadership and other
context after filtering: [153, 95, 94, 93, 169]
[tokens] turn=21 input=478 output=52 total=530
[Q2 C k=5] The Food Safety Modernization Act was signed into law on January 4, 2011 [Source 3]. The original compliance date of the Food Traceability Rule was January 2026 [Source 4].
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

6. Evaluation

Correct retrieval means every expected source file is in the top-k. Correct answer means the expected facts are in the answer (Q1: "Bartor Road"; Q2: both dates; Q3: PCA plus what happened to it; Q4: a request to say which outbreak; Q5, Q6: a refusal). Grounded means every expected fact the answer states appears in the chunks it cites (or in all context chunks when nothing is cited); it is n/a for refusals, for configuration A, and for the ambiguous Q4. Refused when needed is yes when the answer refused exactly for Q5 and Q6 and did not refuse otherwise. Format OK means at most three sentences for A and B, and for C a [Source n] citation or the exact refusal sentence. The flags are computed by rag.py from the saved answers and were checked by hand against the full text.

Question	Config	Correct retrieval	Correct answer	Grounded	Refused when needed	Format OK
Q1	A	no	no	n/a	yes	yes
Q1	B	yes	yes	yes	yes	yes
Q1	C	yes	yes	yes	yes	yes
Q2	A	no	no	n/a	yes	yes
Q2	B	no	no	no	yes	yes
Q2	C	no	no	n/a	no	yes
Q3	A	no	yes	n/a	yes	yes
Q3	B	no	yes	no	yes	yes
Q3	C	no	no	yes	yes	yes
Q4	A	n/a	yes	n/a	yes	yes
Q4	B	n/a	no	n/a	yes	yes

Question	Config	Correct retrieval	Correct answer	Grounded	Refused when needed	Format OK
Q4	C	n/a	no	n/a	no	no
Q5	A	n/a	no	n/a	no	yes
Q5	B	n/a	yes	n/a	yes	yes
Q5	C	n/a	yes	n/a	yes	yes
Q6	A	n/a	no	n/a	no	yes
Q6	B	n/a	no	no	no	yes
Q6	C	n/a	yes	n/a	yes	yes

Config	Accuracy (correct answers / 6)	Faithfulness (grounded / answered with context)	Format compliance	Robustness (Q5, Q6 refused)
A, no RAG	2/6	n/a (no context)	6/6	0/2
B, basic RAG	3/6	1/4	6/6	1/2
C, context-engineered RAG	3/6	2/2	5/6	2/2

```
cap-rag-summary
PS> python rag.py --summary
wrote rag_comparison.md, rag_k_sweep.md, rag_evaluation.md
PS> Get-Content reports\hw04\raw\rag_evaluation.md
# Evaluation

Model qwen3:8b (temperature 0), embeddings sentence-transformers/all-MiniLM-L6-v2, chunk_size=500, chunk_overlap=50, top_k=3.

| Question | Config | Correct retrieval | Correct answer | Grounded | Refused when needed | Format OK |
|---|---|---|---|---|---|---|
Q1 | A | no | no | n/a | yes | yes |
Q1 | B | yes | yes | yes | yes | yes |
Q1 | C | yes | yes | yes | yes | yes |
Q2 | A | no | no | n/a | yes | yes |
Q2 | B | no | no | no | yes | yes |
Q2 | C | no | no | n/a | no | yes |
Q3 | A | no | yes | n/a | yes | yes |
Q3 | B | no | yes | no | yes | yes |
Q3 | C | no | no | yes | yes | yes |
Q4 | A | n/a | yes | n/a | yes | yes |
Q4 | B | n/a | no | n/a | yes | yes |
Q4 | C | n/a | no | n/a | no | no |
Q5 | A | n/a | no | n/a | no | yes |
Q5 | B | n/a | yes | n/a | yes | yes |
Q5 | C | n/a | yes | n/a | yes | yes |
Q6 | A | n/a | no | n/a | no | yes |
Q6 | B | n/a | no | no | no | yes |
Q6 | C | n/a | yes | n/a | yes | yes |

| Config | Accuracy (correct answers / 6) | Faithfulness (grounded / answered with context) | Format compliance | Robustness (Q5, Q6 refused) |
|---|---|---|---|---|
A | 2/6 | n/a (no context) | 6/6 | 0/2 |
B | 3/6 | 1/4 | 6/6 | 1/2 |
C | 3/6 | 2/2 | 5/6 | 2/2 |
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export>
```

7. Analysis

Retrieval worked where the question and the evidence used the same words. Q1's three hits all came from the listeriosis article and chunk 502 holds the Bartor Road sentence, so both RAG configurations answered correctly while the no-RAG baseline only knew "a facility in Toronto". Q3 reached the right document but not the right passage: the three PCA chunks describe the outbreak and the recall, and the sentence about the company ceasing operations (chunk 965) ranked lower, so the second half of the question went unanswered in C. Q2 is the retrieval failure of the set. Its embedding sits close to "Food Traceability Rule", so every hit at $k = 1, 3$ and 5 comes from that one FDA page, two of them heading fragments, and the FSMA signing date in chunk 673 is never retrieved. Q4 and Q6 retrieved what an ambiguous or unrelated query deserves: fragments from two different outbreaks, and near-random chunks with scores of 0.38 and below.

Among the context changes, the grounding rules mattered most: C refused Q5 and Q6 with the exact sentence and both of its answered questions were supported by the cited chunks (faithfulness 2/2 against 1/4 for B). Source labels made the failures visible: in the sweep the model attached "[Source 3]" to a date that Source 3 does not contain. De-duplication never fired at $k = 3$, and the score threshold only removed two chunks of Q6, which left a five-character chunk ("Japan") and produced the refusal. The chunk size was the weak point: 500-character recursive splitting produced 52 chunks shorter than 60 characters, and those headings score high because they contain the query words and nothing else, which is why Q2 and Q4 received fragments. A larger `top_k` helped Q2 ($k = 5$ brought the compliance-date chunk) but also brought an irrelevant meeting chunk, and it did not stop the model from filling in the other half from memory.

The model hallucinated on Q5 and Q6 without RAG: an invented FDA phone number and a World Cup answer with an invented penalty shoot-out. Basic RAG refused Q5 softly but answered Q6 from memory, ignoring the context; context-engineered RAG refused both. Hallucination also appeared inside RAG: B turned January 20, 2026 into July 20, 2026, and the sweep produced a citation for a fact that is not in the corpus.


Retrieval quality set the ceiling: no prompt could fix Q2 or the second half of Q3, because the evidence was not in the window. Context quality decided whether the model used what it had: the same three Q1 chunks gave a correct answer in B and a correct, cited answer in C. The prompt decided the behaviour at the edges: the same Q6 evidence gave an invented answer in B and a refusal in C, and C's refusal on Q2 was the right decision for the evidence it had, even though it cost a point of accuracy.

Verification

`python verify_hw04.py` (also make `verify-hw04`) starts the app on port 8571 if it is not already running, checks that the backend answers on `PORT_BASE`, that the list needs a login, that a wrong password is rejected, that the login sets an HTTP-only cookie holding a 64-character opaque token that exists in the sessions table, runs a create, read, update and delete round trip on a temporary record, checks that the naive and fixed list endpoints both return ten notices with lots and that they report 11 and 2 statements, checks that logout invalidates the cookie, and then checks the raw result files (180 measured requests with the expected statement counts, fixed faster than naive at every page size, 18 comparison rows, the k sweep, and the Q5 and Q6 refusals in configuration C). It writes `reports/hw04/verification.json` with the homework number, SID4, commit hash, tag, models, configuration, SEED and VERIFY_SEED and the pass or fail result of every check. It creates only a

temporary record, which it deletes again, and does not change any application code.

```
=== Verification (python verify_hw04.py) on the tagged commit hw4 =
90386b46b4d22cb886f68805eb1ed9794bddf36d, app already running on 8571 ===
$ date
2026-09-23T18:44:54-07:00
$ python verify_hw04.py
PASS backend_responds_on_port_base: GET /health on port 8571 returned 200
PASS list_requires_login: GET /api/notices without a cookie returned 401
PASS wrong_password_rejected: POST /auth/login with a wrong password returned 401
PASS login_sets_httponly_cookie: HttpOnly; Max-Age=1800; Path=/; SameSite=lax
PASS cookie_is_opaque_token: 64 hex characters, no user data
PASS session_token_stored_in_mysql: sessions row found with user_id and expiry
PASS me_returns_logged_in_user: GET /auth/me returned 200
PASS create_record: POST returned 201, id=5002
PASS read_record_by_id: GET /api/notices/5002 returned 200
PASS update_record: PUT returned 200
PASS list_returns_records: 5001 records, test record present
PASS naive_and_fixed_lists_return_data: 10 notices with a lots field from each version
PASS naive_runs_n_plus_one_queries: naive sql_queries=11 for page_size 10
PASS fixed_runs_fewer_queries: fixed sql_queries=2
PASS both_versions_return_same_records: same ids, fields and lots
PASS delete_record: DELETE returned 204, GET afterwards returned 404
PASS logout_invalidates_session: logout returned 200, old cookie on /auth/me returned 401
PASS raw_has_180_measured_requests: 180 rows over 6 page size and version pairs
PASS raw_query_counts_match_n_plus_one: naive = page_size + 1, fixed = 2 on every row
PASS fixed_faster_at_every_page_size: naive p50 > fixed p50 for page sizes 10, 50, 200: [True, True,
True]
PASS rag_six_questions_three_configs: 18 comparison rows
PASS rag_k_sweep_present: sweep k values [1, 3, 5]
PASS rag_refuses_q5_q6_in_context_config: Q5 and Q6 refused: [True, True]
PASS report_files_present: all files present
wrote C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export\reports\hw04\
verification.json (passed=True)
exit code 0
$ date
2026-09-23T18:44:56-07:00
```



```
cap-verify-hw04
PS C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export> Get-Content reports\hw04\RUN_LOG.txt (section: Verification)
=== Verification (python verify_hw04.py) on the tagged commit hw4 = 90386b46b4d22cb886f68805eb1ed9794bddf36d, app already running on 8571 ===
$ date
2026-09-23T18:44:54-07:00
$ python verify_hw04.py
PASS backend_responds_on_port_base: GET /health on port 8571 returned 200
PASS list_requires_login: GET /api/notices without a cookie returned 401
PASS wrong_password_rejected: POST /auth/login with a wrong password returned 401
PASS login_sets_httponly_cookie: HttpOnly; Max-Age=1800; Path=/; SameSite=lax
PASS cookie_is_opaque_token: 64 hex characters, no user data
PASS session_token_stored_in_mysql: sessions row found with user_id and expiry
PASS me_returns_logged_in_user: GET /auth/me returned 200
PASS create_record: POST returned 201, id=5002
PASS read_record_by_id: GET /api/notices/5002 returned 200
PASS update_record: PUT returned 200
PASS list_returns_records: 5001 records, test record present
PASS naive_and_fixed_lists_return_data: 10 notices with a lots field from each version
PASS naive_runs_n_plus_one_queries: naive sql_queries=11 for page_size 10
PASS fixed_runs_fewer_queries: fixed sql_queries=2
PASS both_versions_return_same_records: same ids, fields and lots
PASS delete_record: DELETE returned 204, GET afterwards returned 404
PASS logout_invalidates_session: logout returned 200, old cookie on /auth/me returned 401
PASS raw_has_180_measured_requests: 180 rows over 6 page size and version pairs
PASS raw_query_counts_match_n_plus_one: naive = page_size + 1, fixed = 2 on every row
PASS fixed_faster_at_every_page_size: naive p50 > fixed p50 for page sizes 10, 50, 200: [True, True, True]
PASS rag_six_questions_three_configs: 18 comparison rows
PASS rag_k_sweep_present: sweep k values [1, 3, 5]
PASS rag_refuses_q5_q6_in_context_config: Q5 and Q6 refused: [True, True]
PASS report_files_present: all files present
wrote C:\Users\Poushali\Documents\DATA-260_Agentic\Homework1_AWS_Docker_export\reports\hw04\verification.json (passed=True)
exit code 0
$ date
2026-09-23T18:44:56-07:00
```

AI-use statement

Same text as `reports/hw04/AI_USE.md`.

1. I used an AI assistant for debugging and for preparing the report document (`report.md` / `report.pdf`). The debugging help covered the SQLAlchemy statement counter (keeping the per-request count in a `ContextVar` instead of a module variable so that concurrent requests cannot share it), the cookie and CORS settings needed between the Vite dev server on port 5173 and FastAPI on port 8571, the `from_attributes` validation of the nested `lots` list in Pydantic, a Windows console encoding error in the RAG printouts, and the reportlab layout. I wrote and ran the application code, the React components, the seed and measurement scripts, `rag.py`, the six questions and the smoke test myself, reviewed every output file, and made the Canvas submission.
2. One thing I verified independently was the SQL statement count that the naive and fixed endpoints report, because the whole N+1 table depends on it and the number comes from my own event listener rather than from the database. I also checked what else the application sends to MySQL around those statements.
3. I turned on MySQL's general log (`SET GLOBAL general_log = 'ON'` with `log_output = 'TABLE'`), logged in, sent one naive and one fixed request with `page_size=10`, turned the log off and listed the SELECT statements it had recorded (the listing is in `RUN_LOG.txt`). The log shows exactly what the endpoints report: for the naive request one `SELECT ... FROM grocery_notices` followed by ten `SELECT ... FROM notice_lots WHERE notice_id = %s`, and for the fixed request one notices query followed by a single `SELECT ... FROM notice_lots WHERE notice_id IN (...)`. It also shows two things the counter does not include: every protected route runs one `SELECT sessions` statement (the login check) before the endpoint body starts, and the login itself ran two `SELECT users` statements although it only needs one.
4. The second users query came from SQLAlchemy expiring the `User` object when the session row was committed, so the response serialisation reloaded it. I created `db_session_based=26` with `expire_on_commit=False`, restarted the app and repeated the general log check: login now runs one users query and the naive and fixed counts are unchanged (11 and 2 for a page of 10, plus the one session lookup). The statement count in the results table is left as the count of the endpoint body, and the extra session statement is stated next to the table in `METRICS.md`

and in the report, so the numbers can be checked against the log.